

---

# Inference as Consolidation: Continual Learning with No Offline Phase by Replay in the Silent Degrees of Freedom

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 An agent that keeps learning in deployment may not be able to afford consolidation  
2 downtime, yet replay-based continual learning almost always rehearses in an offline  
3 phase or interleaves replayed samples with the incoming stream. Brains face the  
4 same constraint and solve part of it with *local sleep*: brief, use-dependent off-  
5 periods of individual circuits during wakefulness. We ask whether a network  
6 trained by local, biologically constrained rules can consolidate *during inference*,  
7 with no offline phase at all. We introduce (i) an *isolation* rule that confines replay  
8 updates to hidden synapses invisible to the current input under  $k$ -winner-take-all  
9 dynamics, exact for silent and suppressed units when the mask matches the weights  
10 being updated, near-exact elsewhere and under the implemented one-step mask  
11 lag by measurement; (ii) a *refractory rotation* rule under which units that have  
12 just fired sit out the next competition, which widens the isolated channel; and (iii)  
13 internal control signals, a homeostatic pressure that triggers replay bursts and a  
14 relative-novelty gate on rotation. The construction inverts the usual direction of  
15 non-interfering continual learning: the hidden computation on the current input  
16 is held invariant while past memories are consolidated in the degrees of freedom  
17 the current batch leaves unused. On class-incremental split-MNIST the system  
18 reaches  $91.6 \pm 0.3\%$  with no offline phase, at or above the best offline-rehearsal  
19 schedule on two held-out splits and consistently above BP+ER and unmasked  
20 local replay, at about twice the offline schedule’s replay samples; in a single pass  
21 over the stream the lead widens (91.8 against 88.7 for BP+ER and 75.1 for offline  
22 rehearsal). Rotation carries most of the gain; isolation adds the guarantee and, at  
23 two-sample replay batches, avoids the severe failures that unmasked replay shows  
24 in half the seeds. The ordering transfers to split CIFAR-10. A drive-referenced  
25 controller for the rule’s synaptic decay replaces per-dataset tuning with one ratio  
26 target and, with locally learned skip synapses, holds a five-hidden-layer network at  
27 the two-layer level where the tuned fixed decay collapses.

## 28 1 Introduction

29 Foundation-model and embodied agents are increasingly expected to keep learning after deployment,  
30 from streams that are non-stationary and never task-labelled. The dominant tool against forgetting  
31 is replay from an episodic memory, and nearly every replay-based method consumes it in one of  
32 two ways: in a dedicated offline phase, a “night” during which the agent stops acting and rehearses,  
33 or by interleaving replayed samples with the live stream. Both cost an agent. An offline phase is  
34 downtime during which the system is unavailable or silently changing; interleaved replay perturbs

the computation currently serving the input, and grows fragile as replay batches shrink toward what an online budget allows.

Biology suggests a third option. Sleep-like off-periods occur *locally*, in individual cortical circuits of awake animals [Vyazovskiy et al., 2011], inducing such ON/OFF alternation in one hemisphere of awake mice discharges local sleep pressure, bilateral induction during sleep deprivation rescues memory consolidation, and an equal tonic reduction of firing does not discharge the pressure [Driessen et al., 2026]. Some core functions of sleep can thus be discharged locally during wakefulness. We take this as an algorithmic motivation, not as a claim about where biological consolidation occurs, and ask whether the computational degrees of freedom a learner is not using can host consolidation while it acts, and what that costs.

We study the question in a deliberately constrained learner: a two-hidden-layer network trained end to end by local, biologically motivated rules (no weight transport, a bounded burst-like error channel, Dale’s law, sparse population codes, sparse distance-dependent wiring, and a single forward–feedback sweep). These constraints expose exactly which degrees of freedom the current input does and does not use. The contributions are:

- **Isolated replay during inference.** A synapse may take a replay update only if its presynaptic unit is silent for the current input or its postsynaptic unit is asleep. Under  $k$ -WTA dynamics this leaves the hidden computation exactly invariant for silent-presynaptic and suppressed units (Propositions 1–2); we measure how nearly it holds elsewhere. With rotation on, replay micro-batches of 8–16 samples suffice, where unmasked replay without rotation and offline schedules degrade. Isolation adds 0.1–0.3 points there (2.6 without rotation) together with the guarantee, and at two-sample batches it shows none of the severe failures that unmasked replay shows in half the seeds. Replay confined to the readout falls to 42%, so hidden-layer replay updates are essential.
- **Refractory rotation.** Units that fired on the current batch sit out the next competition. This alternation, the algorithmic counterpart of the ON/OFF requirement in Driessen et al. [2026], widens the consolidable channel from  $\approx 0.25$  to 0.53 of the replay-active units and is worth +3.6 sequential points on top of isolation (+6.1 alone). A random alternation of matched size recovers most of that, and use-dependence the remaining 0.6 sequential and 2.0 static points; a tonic reduction, like the biological control, loses on both axes.
- **Internal triggers,** as schedulers for sparse replay budgets (the headline system replays at every batch). Replay is triggered by a unit-level homeostatic pressure; the tested surprise trigger fails during wake, whereas offline rehearsal is best timed by novelty in our sweeps. A relative-novelty gate on rotation is scale-free where absolute thresholds fail across the regimes considered here.
- **One ratio target instead of per-dataset decay tuning, and depth.** A drive-referenced controller for the rule’s synaptic decay removes its last per-dataset constant and, with locally learned skip synapses, holds a five-hidden-layer network at the two-layer level on both axes where the tuned fixed decay without skips collapses to chance.

Everything is evaluated on two axes at once, the class-incremental stream and the identical learner on the i.i.d. shuffle of the same data, because consolidation during inference must not damage ordinary learning. The scale is small (MLPs on MNIST/CIFAR variants); the object of study is the mechanism and its prerequisites (Section 6).

## 2 Related work

**Sleep-inspired consolidation.** The Bazhenov line implements sleep as a global offline phase driven by noise under local plasticity [Tadros et al., 2022, Bazhenov et al., 2026, Golden et al., 2022, Kubo et al., 2025]; wake–sleep frameworks with NREM/REM stages [Sorrenti et al., 2025, Robinson et al., 2022] and systems such as SIESTA [Harun et al., 2023] likewise schedule consolidation offline, and SESLR [Lin et al., 2026] appends a noise-enhanced sleep phase (a frozen-extractor spiking classifier training only on its buffer) to correct online learning’s recency bias. None consolidates in currently unused units during inference, and the bias correction such a phase delivers is what our plastic readout receives continuously from isolated replay.

**Gating and history-dependent competition.** Context-dependent gating activates a subnetwork per task from task identity [Masse et al., 2018]; learned gates [Tilley et al., 2023], dendritic context [Iyer et al., 2022], task-free sparsification [Lässig et al., 2023] and recovered functional masks [McKee et al., 2026] derive the gate from the input. Refractory competition has classical precedents [Kaski and Kohonen, 1994, Maeda and Miyajima, 1999], and spiking continual learners raise firing thresholds with use [Shen et al., 2024] or randomise the top- $k$  selection [Shen et al., 2026] to de-overlap task codes, the wake-side counterpart of our random-alternation control. Our rotation differs in purpose: recently active units are benched for one competition to enlarge a *replay channel*, not to allocate resources across tasks.

**Null-space and parameter isolation.** A mature line protects *previous* tasks: gradient or activation-subspace projections [Farajtabar et al., 2020, Saha et al., 2021], feature-covariance null spaces [Wang et al., 2021],  $k$ -winner sparsity with such projections [Abbasi et al., 2022], frozen subnetworks [Golkar et al., 2019, Mallya and Lazebnik, 2018]. We invert the protected direction: the ongoing computation is held invariant while old-memory replay is written into degrees of freedom the current sparse support exposes as invisible. No bases or task masks are stored, and the “null space” changes with every batch at no cost. Complementary-learning-system methods [Arani et al., 2022, Pham et al., 2021] raise the value of each replayed sample but still replay at every step; learned replay scheduling [Klasson et al., 2023] is related to our trigger study.

### 3 Method

**Setting.** A stream of labelled examples is presented in  $T$  contiguous tasks with disjoint class sets; no task identity is available at test time; a single shared  $C$ -way readout is evaluated on all classes after the last task (class-incremental learning). We report final accuracy  $A_{\text{seq}}$ , forgetting  $F$ , the accuracy  $A_{\text{iid}}$  of the identical learner and budget on the i.i.d. shuffle (the *static* axis), the replay cost  $R$  in samples, and an arithmetic-energy proxy  $E$ ; the learner is confined to the cortical constraint set  $\mathcal{B}$  above and to an episodic memory of  $K$  raw examples with reservoir writing [Vitter, 1985]. The question is whether the offline term of consolidation can be removed entirely without losing  $A_{\text{seq}}$ , and at what price in  $A_{\text{iid}}$  and  $R$ .

**Base learner.** Layer activities are  $a^0 = x$  and, for  $\ell = 1, \dots, L$ ,

$$z^\ell = W^\ell a^{\ell-1} + b^\ell, \quad a^\ell = \Phi_k(\sigma(z^\ell) \odot (1 - s^\ell)), \quad (1)$$

where  $\sigma$  is a rectifier,  $s^\ell \in \{0, 1\}^{n_\ell}$  the suppression mask of the rotation rule below, and  $\Phi_k$  keeps the  $k$  largest entries of each row ( $k$ -WTA); a per-unit homeostatic gain in the code is fixed at 1. Layers  $1, \dots, L-1$  are hidden; layer  $L$  is a linear  $C$ -way readout whose pre-activation  $z^L$  is the output, scored by squared error against the one-hot target  $y$ . Errors come from one top-down sweep through learned feedback weights  $B^\ell \in \mathbb{R}^{n_{\ell+1} \times n_\ell}$  (no weight transport),

$$\varepsilon^L = y - z^L, \quad \varepsilon^\ell = \sigma'(z^\ell) \odot \left( (B^\ell)^\top [\kappa \tanh(\varepsilon^{\ell+1}/\kappa) \odot \mathbf{1}(a^{\ell+1} > 0)] \right), \quad (2)$$

a bounded ( $|\varepsilon_i^\ell| \leq \kappa \sum_j |B_{ji}^\ell|$  per component), event-gated burst signal in the spirit of predictive-coding and burst-dependent learners [Whittington and Bogacz, 2017, Payeur et al., 2021]. Forward and feedback weights follow the same local product with a shared decay  $\lambda$  (Kolen–Pollack alignment [Kolen and Pollack, 1994, Akrouf et al., 2019]),  $\Delta W^\ell \propto \varepsilon^\ell (a^{\ell-1})^\top - \lambda W^\ell$ ,  $\Delta B^{\ell-1} \propto \varepsilon^\ell (a^{\ell-1})^\top - \lambda B^{\ell-1}$ , under Dale’s law with sign-concordant feedback. The optimiser is heavy-ball SGD, with one per-synapse velocity and no second-moment state (Section 4: masked Adam ties it, and momentum-free SGD does not serve inside the consolidation loop). The base learner costs 1.2 points against a dense backpropagation MLP on i.i.d. MNIST at  $4\times$  fewer synaptic operations (Appendix C).

**Isolated replay.** Let  $X = \{x_b\}_{b=1}^m$  be the current waking batch,  $a_{i,b}^\ell$  the activity of unit  $i$  on sample  $b$ , and  $\mathcal{A}^\ell(X) = \{i : \exists b, a_{i,b}^\ell \neq 0\}$  the awake set. During the same forward step, a replay micro-batch drawn from the buffer is applied through the synapse mask

$$M_{ij}^\ell = \mathbf{1}[j \notin \mathcal{A}^{\ell-1}(X) \vee i \notin \mathcal{A}^\ell(X)], \quad (3)$$

i.e. iff the presynaptic unit is silent for the current input or the postsynaptic unit is asleep. The mask applies to the *entire* replay-induced parameter change, not only to the data gradient: with heavy-ball

132 velocity  $v$  ( $\mu = 0.9$ ), the local update direction  $G = \varepsilon^\ell(a^{\ell-1})^\top$  and the per-step decay coefficient  $\lambda$ ,  
 133 a replay step is

$$v \leftarrow M \odot (\mu v + G) + (1 - M) \odot v, \quad W \leftarrow W + \eta M \odot v - \lambda M \odot W, \quad (4)$$

134 so the velocity is frozen (not zeroed) outside the mask, the decay is confined to it, and biases follow  
 135 the unit mask  $\mathbf{1}[i \notin \mathcal{A}^\ell(X)]$ ; unmasked momentum would re-inject the update into awake synapses  
 136 [McKee et al., 2026, Bricken et al., 2023]. The readout row and its bias stay plastic, so old classes  
 137 remain calibrated; the guarantees below concern the hidden computation. Within one step the order  
 138 is: forward pass under the current suppression, unmasked waking update, awake sets from that pass,  
 139 masked replay on the updated weights; replay is inferred with the suppression removed (a refractory  
 140 unit may fire on replay and learn from it) and a burst reuses one mask (Algorithm 1). Throughout,  
 141 “during inference” means between the steps of the stream: the three operations run in sequence within  
 142 each step, and we claim nothing about asynchronous execution or serving latency.

143 **Proposition 1** (Pre-silent updates are exactly invisible). *Fix  $X$  and let  $\widehat{W}^\ell = W^\ell + \Delta^\ell$  with  $\Delta_{ij}^\ell = 0$   
 144 whenever  $j \in \mathcal{A}^{\ell-1}(X)$ . Then every hidden activity on every sample of  $X$  is unchanged for any  
 145 magnitude of  $\Delta$ ; if the readout row satisfies the same condition, or is held fixed, the network output is  
 146 unchanged as well.*

147 **Proposition 2** (Post-asleep updates are invisible under a margin). *Let  $\Delta^\ell$  and  $\Delta b^\ell$  additionally  
 148 move synapses and biases of units  $i \notin \mathcal{A}^\ell(X)$  with active presynaptic partners, and write  $\delta_{i,b} =$   
 149  $\sum_j \Delta_{ij}^\ell a_{j,b}^{\ell-1} + \Delta b_i^\ell$ . If for every such unit and every sample  $\sigma(z_{i,b}^\ell + \delta_{i,b}) < \tau_{k,b}^\ell$ , where  $\tau_{k,b}^\ell$  is the  
 150  $k$ -th largest value entering  $\Phi_k$  on sample  $b$  (strict, so a tie cannot promote the unit; when fewer than  
 151  $k$  entries are positive,  $\tau_{k,b}^\ell = 0$  and the condition reads  $\sigma(\cdot) = 0$ ), the hidden activities on  $X$  are  
 152 unchanged.*

153 **Corollary 1** (Suppressed units are exact for any magnitude). *If  $s_i^\ell = 1$  in (1),  $a_{i,b}^\ell = 0$  for every  
 154 sample and Proposition 2 holds with no margin condition.*

155 Proofs are in Appendix A. The propositions take the mask and the weights it protects from the same  
 156 state; in the implemented order the awake sets precede the waking update, so they hold exactly for the  
 157 pre-update weights and the one-step lag is part of the measured drift (0.02 of the 0.32 hidden-code  
 158 points, Appendix I). The default system does not enforce the margin, so we measured it over all  
 159 16,885 isolated replay updates of the reported run (held-out protocol, three seeds). The update  
 160 altered the top hidden code of 0.33% of waking samples (0.31–0.35) and a waking prediction for  
 161 0.32%, almost all through the plastic readout (0.001% with the readout isolated, one seed); it violated  
 162 the margin for  $\sim 10^{-4}$  of asleep units per step. The isolation is thus exact for silent-presynaptic  
 163 and suppressed units and near-exact elsewhere for the hidden code of the *current* batch under its  
 164 training-time suppression mask (evaluation runs with the mask off). What replay does to later inputs  
 165 is measured by the ablation (Table 1), and Corollary 1 is why rotation makes the isolation robust.

166 **Refractory rotation.** After batch  $X_t$ , set  $s_i^\ell(t+1) = \mathbf{1}[i \in \mathcal{A}^\ell(X_t)]$ : every unit that fired sits out  
 167 the next competition and is therefore asleep, and exactly consolidable, for  $X_{t+1}$ . Rotation widens  
 168 the isolated channel  $\rho_\ell = |\tilde{\mathcal{A}}^\ell \setminus \mathcal{A}^\ell(X)|/|\tilde{\mathcal{A}}^\ell|$ , the fraction of the units  $\tilde{\mathcal{A}}^\ell$  activated by the replay  
 169 micro-batch that are asleep for the current input, from 0.2–0.3 to  $\approx 0.53$ , at the price of running  
 170 inference on the second-best coalition.

171 **Internal control (Figure 1b).** *When to replay:* each unit integrates its own use,  $S_i \leftarrow S_i + \bar{a}_i(X_t)$   
 172 with  $\bar{a}_i(X_t)$  the fraction of the batch’s samples on which it fired (a unit-level Process S [Borbély,  
 173 1982]), and a burst of  $n_b$  replay micro-batches fires when the mean pressure over the currently asleep  
 174 units of all hidden layers exceeds  $\theta$ ; consolidation resets the pressure of the units that received it  
 175 ( $S_i \leftarrow 0$ ). *When to rotate:* with fast and slow exponential averages of the batch-mean squared output  
 176 error  $e_t$ ,  $\bar{e}_f$  ( $\alpha_f=0.1$ ) and  $\bar{e}_s$  ( $\alpha_s=0.002$ ), rotation runs while  $\bar{e}_f \geq \beta \bar{e}_s$  ( $\beta \approx 1.1$ –1.5) or while  
 177  $\bar{e}_f \geq \gamma e_0$  with  $e_0$  the chance-level error of the first 20 batches ( $\gamma \approx 0.5$ ): the stream is new relative  
 178 to the learner’s own recent history, or is not yet mastered. The detector is a standard drift monitor  
 179 [Ross et al., 2012], loosely inspired by cholinergic novelty signalling [Hasselmo, 2006]; what is new  
 180 is what it gates. A single absolute threshold cannot serve across the regimes considered here, because  
 181 a converged ten-class i.i.d. error exceeds any level a two-class task dips under. At the activity fraction  
 182 used throughout the gate is nearly neutral, and the default system runs rotation always on (Section 4).




Figure 1: **Mechanism.** (a) During each waking batch, replay from the episodic buffer is written only into hidden synapses invisible to the current input, exactly so for silent presynaptic and suppressed units; refractory rotation forces recently used units into the consolidable set. (b) Unit-level homeostatic pressure triggers replay bursts, and a relative-novelty gate decides when rotation runs. The default system runs rotation always on with one micro-batch per waking batch; the pressure trigger and the novelty gate are schedulers for sparse budgets and lower activity fractions.

**Protocol.** Split-MNIST (five tasks of two classes, five epochs per task) and split CIFAR-10 (same protocol,  $32 \times 32$  luminance), each with its i.i.d. counterpart; a third regime feeds CIFAR-10 through a frozen V1-like patch front-end (1024-D, linear probe  $\approx 49\%$ ); split CIFAR-100 appears in Section 5 for mechanism only. Architecture  $d=512-256-C$ ,  $k$ -WTA 10%, 30% distance-dependent connectivity, 80% excitatory units; waking batches of 16; buffer  $K=1000$  ( $K=200$  in Appendix G). Baselines: BP with experience replay (BP+ER) at equal  $K$ ; the offline “night” (20 replay batches of 256 at  $3 \times$  learning rate after each waking epoch, the best offline schedule in our sweep) on its preferred narrower core and on our substrate; interleaved ER for the local learner; no buffer. Three seeds (six on headline rows), mean  $\pm$  s.d. Configurations were selected on the official test splits, which served as the development set, and then frozen. Every main-text comparison re-trains them with a stratified tenth of the training data removed from the stream and evaluates on that tenth, which took part in no selection; a second tenth, drawn with a different seed and sharing 9% of its samples with the first, replicates the headline rows (Appendix L). Results not re-run under this protocol are labelled development-phase. Forgetting  $F$  is the mean drop of each earlier task’s accuracy from its own end to the end of training (per-task matrix in Appendix L). Evaluation runs with the suppression mask off ( $k$ -WTA only), so test order and batch size cannot affect it. Energy is an idealised 45-nm FP32 arithmetic proxy [Horowitz, 2014] (0.9 pJ per accumulate, one per synaptic event, against 4.6 pJ per multiply-accumulate) after rate-to-spike conversion with training-calibrated thresholds [Rueckauer et al., 2017], excluding data movement and membrane updates. Further protocol details are in Appendix D.

## 4 Results

**Local sleep replaces the night.** On split-MNIST at  $K=1000$ , refractory local sleep with one isolated replay micro-batch beside every waking batch reaches  $91.6 \pm 0.3\%$  (six seeds) with no offline phase. This is above the best offline night ( $89.2 \pm 1.7\%$  on its preferred narrow substrate;  $89.0 \pm 3.5$ , highly variable across seeds, on ours; paired margins  $+2.5$  [ $+1.3, +4.8$ ] and  $+2.6$  [ $+0.2, +5.2$ ]), above unmasked interleaved ER at the same budget ( $87.5 \pm 0.8$  under Adam,  $85.4 \pm 3.5$  under SGD) and above BP+ER ( $88.8 \pm 0.3$ ;  $+2.8$  [ $+2.5, +3.2$ ]). In a single pass over the stream (one epoch per task) the ordering sharpens:  $91.8 \pm 0.2$  against  $88.7 \pm 0.3$  for BP+ER and  $75.1 \pm 3.5$  for the offline night, for which one night per task is too few (Table 1). The headline configuration replays  $2.2 \times$  the night’s samples; at  $1.1 \times$  (10% more) the system still reaches  $91.1$ – $91.2$  (eight-sample micro-batches every batch or 16-sample ones every second batch,  $35 \times$  and  $17.6 \times$  the night’s updates; Appendix D). On a second held-out split the ordering holds (rotation  $91.8 \pm 0.4$ , BP+ER  $88.5 \pm 0.7$ , unmasked ER  $80.0 \pm 14.1$ , narrow night  $89.3 \pm 0.6$ ), with one exception: the same-substrate night, bimodal on the first split, reaches  $91.5 \pm 0.5$  on the second and trails by only 0.3 there (Appendix L).

**What rotation and isolation each buy.** Figure 2 (left) varies the replay batch size. The full system is flat from 256 down to 8 ( $91.6 \pm 0.5 \rightarrow 91.1 \pm 0.4$ ) and loses 1.9 at 4 ( $89.7 \pm 0.7$ ). At batch 16, unmasked replay without rotation reaches  $85.4 \pm 3.5$ , isolation without rotation  $88.0 \pm 0.7$ , and the offline night falls from  $89.0$  to  $85.8 \pm 3.0$  when its batches shrink to 16. The fourth cell, rotation with unmasked replay, reaches  $91.5 \pm 0.3$  at batch 16,  $90.9 \pm 0.1$  at 8 and  $89.3 \pm 0.7$  at 4. The factors





Figure 2: **Left (isolation):** split-MNIST accuracy against replay batch size, the number of updates fixed within each series (one per waking batch; 20 per epoch for the night, whose samples therefore fall with its batch): the full system is flat down to 8 samples and loses 1.9 at 4 (12.7 at 2, Appendix K); rotation with unmasked replay (dashed) trails it by 0.1 to 0.3, within seed noise, and fails severely in half the seeds at batch 2; without rotation, unmasked and isolated replay and the offline night all degrade. **Right (timing):** accuracy against the number of waking replay events: trickles and clocked bursts fail alike, pressure-triggered bursts (star) keep most of the accuracy at a third of the events, surprise triggering fires no events in this regime.

are therefore strongly sub-additive: rotation is the larger one (+6.1 alone, +3.6 on top of isolation), and isolation’s seed-paired margin is 0.1  $[-0.3, +0.5]$  at batch 16, 0.3  $[-0.2, +0.6]$  at 8 and 0.3  $[-0.5, +1.2]$  at 4, within seed noise at every size, with 0.4 on the static axis (94.1 vs. 93.7). At two-sample batches both lose heavily, but isolated replay degrades gracefully ( $79.0 \pm 3.2$ , every seed within 75–84) where unmasked replay fails severely (below 60%) in three of six seeds (Appendix K). In the seeds and configurations tested, isolation’s accuracy value is thus the absence of those failures, not a margin. What it buys at every batch size is the guarantee; without it the outcome rests on the rotation’s tolerance of the perturbation.

**Component prices (Table 1).** Removing rotation costs 3.6 sequential points ( $88.0 \pm 0.7$ : the same mask on natural silence alone), removing both rotation and isolation 6.2 with a twelve-fold larger seed deviation ( $85.4 \pm 3.5$ ), removing isolation alone 0.1 (above), and removing replay everything (19.4, the forgetting floor). Replay confined to the readout row falls to  $41.7 \pm 10.9$  and costs 11 static points, so hidden-layer replay updates are essential. A random suppression of matched size recovers most of rotation’s gain ( $91.0 \pm 0.8$ ) but trails it by 0.6 sequential and 2.0 static points with  $1.4\times$  the forgetting: alternating the coalition carries most of the effect, and resting the recent winners the rest. Rotation is the only statically costly component, at 1.2–1.7 points for this activity fraction (94.1 against 95.3–95.8 without it; isolation adds 0.4). The unmasked control is not handicapped by the replay step ( $3\eta$  against a waking step of  $\eta/16$ ): over replay gains  $\{1/16, 1/4, 1, 3, 10\}$  masked replay peaks at the default (91.6) while unmasked replay plateaus from gain 1 to 3 (85.7, 85.4), gain 10 destabilises both, and the rotation-free unmasked control never closes the gap (the sweep bounds that control; Appendix I). *Mirror isolation* (waking updates barred from synapses that served the last replay batch, replay unmasked) reaches only  $78.0 \pm 3.2$ . This reversed-mask control loses fourteen points; it is one implementation of the protect-the-past direction, not a test of the methods in that literature. *Soft rotation* (halving rather than silencing recent winners) loses on both axes ( $89.9 \pm 0.7/84.3 \pm 1.0$ ): all-or-none OFF periods beat turned-down gain. *Optimiser state:* masked Adam and heavy-ball SGD tie on this split ( $91.5 \pm 0.4/94.1 \pm 0.1$  vs.  $91.6 \pm 0.3/94.1 \pm 0.2$ ), and letting Adam’s moments leak outside the mask adds nothing ( $91.6 \pm 0.6$ ); we keep the single velocity as the least optimiser state that stays confined to the mask, not for accuracy. Momentum-free SGD trails by 3.2/4.7 even after its replay step is re-scaled by hand (otherwise it lands at chance), and a two-sided learning-rate sweep shows that these comparisons are plateaus (Appendices I, H).

**When replay must be rare: homeostatic bursts, not surprise.** Figure 2 (right): at 12.5% of the replay events a clocked burst of 16 batches every 128 reaches  $80.5 \pm 6.1$ , no better than an even trickle at the same budget ( $78.5 \pm 2.6$ ); unit-level pressure triggering keeps  $88.6 \pm 1.6$  at a third of the events, and surprise triggering fires no event (18.8). This is the opposite of the offline night, whose best trigger in our sweeps is novelty-timed.

Table 1: **Component ablation** (512–256,  $k$ -WTA 10%,  $K=1000$ , one replay micro-batch of 16 per waking batch; held-out split). Sequential = split-MNIST; static = i.i.d. MNIST; single pass = one epoch per task, otherwise identical. Headline rows six seeds, the rest three.

| Variant                                                                   | Sequential                       | i.i.d. (static) | Seq., single pass |
|---------------------------------------------------------------------------|----------------------------------|-----------------|-------------------|
| <b>full system:</b> isolated replay + always-on rotation (heavy-ball SGD) | <b>91.6 <math>\pm</math> 0.3</b> | 94.1 $\pm$ 0.2  | 91.8 $\pm$ 0.2    |
| rotation gated by relative novelty, committed to bouts                    | 91.8 $\pm$ 0.3                   | 94.4 $\pm$ 0.3  |                   |
| rotation gated per batch                                                  | 87.5 $\pm$ 1.6                   | 94.7 $\pm$ 0.1  |                   |
| – rotation (same mask, Eq. 3, natural silence only)                       | 88.0 $\pm$ 0.7                   | 95.3 $\pm$ 0.2  |                   |
| – isolation, – rotation (unmasked interleaved replay)                     | 85.4 $\pm$ 3.5                   | 95.8 $\pm$ 0.3  | 76.0 $\pm$ 20.1   |
| – isolation, + rotation (rotation alone)                                  | 91.5 $\pm$ 0.3                   | 93.7 $\pm$ 0.3  |                   |
| readout-only replay (hidden layers take no replay update)                 | 41.7 $\pm$ 10.9                  | 82.7 $\pm$ 0.3  |                   |
| random rotation (matched suppression count)                               | 91.0 $\pm$ 0.8                   | 92.1 $\pm$ 0.2  |                   |
| – replay (no buffer)                                                      | 19.4 $\pm$ 0.0                   | 95.1 $\pm$ 0.0  | 18.1 $\pm$ 0.7    |
| masked Adam, moments confined to the mask                                 | 91.5 $\pm$ 0.4                   | 94.1 $\pm$ 0.1  |                   |
| masked Adam, moments leaking outside the mask                             | 91.6 $\pm$ 0.6                   | 94.0 $\pm$ 0.3  |                   |
| momentum 0 (pure SGD; $\eta$ and replay gain re-tuned, else chance)       | 88.4 $\pm$ 0.5                   | 89.4 $\pm$ 0.3  |                   |
| offline night instead of concurrent replay, same substrate                | 89.0 $\pm$ 3.5                   | —               | 75.1 $\pm$ 3.5    |
| the same night on its preferred narrow substrate                          | 89.2 $\pm$ 1.7                   | —               |                   |
| mirror direction: protect the past                                        | 78.0 $\pm$ 3.2                   | 91.7 $\pm$ 0.6  |                   |
| soft rotation: gain 0.5 instead of 0                                      | 89.9 $\pm$ 0.7                   | 84.3 $\pm$ 1.0  |                   |
| unmasked ER at the same budget (Adam)                                     | 87.5 $\pm$ 0.8                   | 95.1 $\pm$ 0.2  |                   |
| backprop (+ER sequential; plain static), dense                            | 88.8 $\pm$ 0.3                   | 97.4 $\pm$ 0.1  | 88.7 $\pm$ 0.3    |

**The i.i.d. price and the activity fraction.** Rotation’s static price is governed by the substrate’s activity fraction: 1.5 points at 5% (92.6 against 94.1 rotation-free) and 1.2–1.7 at 10%, where the gates are neutral (Table 1). From 10 to 25% the sequential axis is flat (91.6, 91.4, 91.5) while the static axis rises (94.1, 94.4, 94.6). On feature CIFAR-10, relaxing the fraction from 5 to 25% improves both axes monotonically (40.9  $\rightarrow$  42.9 / 38.3  $\rightarrow$  46.4), with the local learner 1.4–1.9 points ahead of matched-sparsity BP at every level (Appendices E–F).

**Transfer to CIFAR-10.** On raw split CIFAR-10 the ordering transfers and widens: refractory local sleep  $28.9 \pm 1.0 >$  night  $25.7 \pm 1.2 >$  BP+ER  $24.8 \pm 0.3 >$  unmasked ER  $15.4 \pm 4.2 \approx$  no buffer 16.3 (the same on the second split:  $29.7 > 26.5 > 25.8 > 15.5$ ). On the feature front-end the local ordering persists (refractory  $42.3 \pm 0.6 >$  night  $34.0 \pm 0.8 \gg$  unmasked  $20.5 \pm 1.0$ ), but BP+ER reaches  $42.8 \pm 0.7$  at its best learning rate, a 0.5-point lead. Transplanting the local network’s ingredients into the BP net one at a time attributes that gap to activation sparsity (+ $k$ -WTA  $40.4 \pm 1.5$ , below the local learner on the identical substrate; Figure 6). Converted to spiking inference ( $T_s=8$ ) the default system runs at 92.7% for a proxy of 117 nJ per sample,  $21\times$  below dense rate inference (Appendix D).

## 5 Scaling the mechanism: self-referenced decay and depth

**The last tuned constant.** Every result above runs the rule with a decay  $\lambda$  tuned once per dataset ( $10^{-3}$  on the image benchmarks). Split CIFAR-100 exposes what that constant hides. Each class receives a tenth of the supervision events while the decay acts on every synapse at every step, so at the MNIST-tuned value the hidden weights follow the pure-decay trajectory and the network lands at chance (2.4% i.i.d. against 20–23% on a ten-class subset, development-phase numbers). We replace the constant by a per-layer ratio controller,

$$\lambda_\ell = \min\left(\lambda_{\max}, \rho \text{EMA}[u^\ell] / (\overline{|W^\ell|} + \epsilon)\right), \quad \rho = 0.25, \lambda_{\max} = 0.02, \quad (5)$$

where  $u^\ell$  is the mean magnitude of the post-optimiser, *pre-decay* waking increment of layer  $\ell$  (the decay never feeds its own drive; EMA coefficient 0.02 over unmasked waking steps,  $\epsilon = 10^{-12}$ ), and  $W^\ell$ ,  $B^{\ell-1}$  share  $\lambda_\ell$ : the decay’s mean magnitude  $\lambda_\ell \overline{|W^\ell|}$  is held at the fraction  $\rho$  of the mean data-driven update.  $\rho$  is the only control target;  $\lambda_{\max}$ , the EMA coefficient and  $\epsilon$  are global safeguards never tuned per dataset. Isolation and rotation decide *where* and *when* to consolidate, the controller *how strongly*.




Figure 3: **Depth is the same starvation on another axis, and two mechanisms cross it.** Sequential (left) and static (middle) accuracy against hidden depth (three seeds; six at the skip system’s depth five): the fixed decay without skips collapses, the controller degrades gracefully, controller + skips hold the two-layer level. Right: the controller’s realised per-layer decay.

285 **One ratio target, every dataset.** With the same  $\rho$  on every dataset (Table 4 in Appendix J; e.g.  
 286 split-MNIST  $91.6/94.1 \rightarrow 90.6/93.5$ , CIFAR-10  $28.9/28.0 \rightarrow 26.2/30.1$ , sequential/static) the  
 287 controller concedes 1–3 sequential points to the per-dataset tuned value and gains 1.2–5.3 static  
 288 points on CIFAR-10, feature CIFAR-10 and both CIFAR-100 variants. On MNIST it gains 1.1 on  
 289 the second held-out split ( $95.2 \pm 0.4$ ) but one of six seeds collapsed on the first ( $93.5 \pm 4.3$ ), so we  
 290 claim no static gain there. On CIFAR-100 it lifts the system  $2.7$ – $4.7\times$  over the manual rescue. The  
 291 realised  $\lambda_\ell$  spans  $10^{-5}$  (MNIST streams) to  $3 \times 10^{-3}$  (raw CIFAR), so the per-dataset ladder was a  
 292 signal-to-decay ratio in disguise (Appendix J).

293 **Depth (Figure 3).** Under the tuned fixed decay without skips the static axis dies at four hidden  
 294 layers ( $10.1 \pm 0.2\%$ , chance) and the sequential axis collapses at five ( $26.3 \pm 20.2$ ); under the  
 295 controller every cell stays alive ( $90.7 \rightarrow 88.9 \rightarrow 88.1$  sequential,  $95.3 \rightarrow 94.8 \rightarrow 89.6$  static). The  
 296 residual toll is architectural. ResNet-style skip synapses (Appendix B) restore the thick error path,  
 297 and the five-hidden-layer system reaches  $91.0 \pm 0.3/95.8 \pm 0.3$  (six seeds) against the two-layer  
 298 default’s  $91.6/94.1$  and the two-layer controller’s  $90.6/93.5$ : 0.6 sequential points below the default  
 299 and 1.7 above it on the static axis. The bypass alone also rescues the fixed decay ( $90.4/93.5$  at  
 300 depth five), so error-path attenuation is a principal cause of the collapse (the fixed decay was not  
 301 re-tuned per depth). BP+ER at the same widths is depth-insensitive ( $88.9 \rightarrow 87.9$ ) but forgets  $2\times$   
 302 more. On CIFAR-100, by contrast, depth is a net loss that skips only half refund, while width pays  
 303 (Appendix J).

## 304 6 Discussion

305 **What this offers agents.** The construction is a recipe. If an agent’s representation is sparse enough  
 306 that a batch leaves most units silent, and its learning rule yields a local error signal, consolidation  
 307 can run in the silent degrees of freedom between the steps of the stream, with no separate offline  
 308 phase, under a conditional guarantee on the computation in use, and at a higher total compute ( $6\times$  the  
 309 night’s wall-clock, unoptimised). Both prerequisites can be checked on a given architecture. Top- $k$   
 310 routed mixture-of-experts layers expose, per token, the experts a token leaves silent, but whether the  
 311 routes of a whole batch leave enough experts idle is a measurement (load balancing pushes against  
 312 it), and the local error signal is what such a system would still have to supply. We treat the biology as  
 313 analogy, not as mechanistic equivalence; the refractory rule mirrors the ON/OFF finding of Driessen  
 314 et al. [2026], and the trigger dissociation is a testable prediction.

315 **Limitations.** All results are at MLP scale on MNIST/CIFAR variants with three to six seeds. Every  
 316 main comparison is on held-out training data (two splits for the headline rows, with exploratory  
 317 results labelled), and the protocol is five epochs per task (a single pass keeps the lead, Table 1).  
 318 The lead over the best offline night is 2.5 points on the first split and 0.3 on the second, where the  
 319 night’s seed variance vanished; the leads over BP+ER and unmasked replay hold on both. BP+ER  
 320 is an external reference rather than a state-of-the-art claim (no DER++-style baselines) and keeps  
 321 a 0.5-point lead on feature CIFAR-10. Rotation keeps a static price of 1–2 points, CIFAR-100 is  
 322 at the floor for every method, and the energy figure is an arithmetic proxy. Within these limits  
 323 the conclusion stands: consolidation does not require a separate offline phase, at  $2.2\times$  the night’s  
 324 replay samples and, on the first split, at  $1.1\times$ . Code can be found at <https://anonymous.4open.science/r/IaC-CL-with-no-offline>  
 325



## References

- Abbasi, A., Nooralinejad, P., Braverman, V., Pirsiavash, H., Kolouri, S. (2022). Sparsity and heterogeneous dropout for continual learning in the null space of neural activations. In *Conference on Lifelong Learning Agents (CoLLAs)*, PMLR 199.
- Akrout, M., Wilson, C., Humphreys, P., Lillicrap, T., Tweed, D. (2019). Deep learning without weight transport. In *NeurIPS* 32.
- Arani, E., Sarfraz, F., Zonooz, B. (2022). Learning fast, learning slow: A general continual learning method based on complementary learning system. In *ICLR*.
- Bazhenov, A., Delanois, J.E., Krishnan, G.P. (2026). Not just after one: Sleep-inspired replay prevents catastrophic forgetting after sequential tasks. *arXiv:2606.08447*.
- Benna, M.K., Fusi, S. (2016). Computational principles of synaptic memory consolidation. *Nature Neuroscience*, 19:1697–1706.
- Borbély, A.A. (1982). A two process model of sleep regulation. *Human Neurobiology*, 1:195–204.
- Bricken, T., Davies, X., Singh, D., Krotov, D., Kreiman, G. (2023). Sparse distributed memory is a continual learner. In *ICLR*.
- Driessen, K., Squarcio, F., Tononi, G., Cirelli, C. (2026). Induction of cortical ON/OFF periods in awake mice fulfills sleep functions. *Nature Neuroscience*, 29:1954–1965.
- Farajtabar, M., Azizan, N., Mott, A., Li, A. (2020). Orthogonal gradient descent for continual learning. In *AISTATS*, PMLR 108.
- Golden, R., Delanois, J.E., Sanda, P., Bazhenov, M. (2022). Sleep prevents catastrophic forgetting in spiking neural networks by forming a joint synaptic weight representation. *PLoS Computational Biology*, 18:e1010628.
- Golkar, S., Kagan, M., Cho, K. (2019). Continual learning via neural pruning. *arXiv:1903.04476*.
- Harun, M.Y., Gallardo, J., Hayes, T.L., Kemker, R., Kanan, C. (2023). SIESTA: Efficient online continual learning with sleep. *TMLR*.
- Hasselmo, M.E. (2006). The role of acetylcholine in learning and memory. *Current Opinion in Neurobiology*, 16:710–715.
- Horowitz, M. (2014). Computing’s energy problem (and what we can do about it). In *IEEE International Solid-State Circuits Conference (ISSCC)*, pp. 10–14.
- Iyer, A., Grewal, K., Velu, A., Souza, L.O., Forest, J., Ahmad, S. (2022). Avoiding catastrophe: Active dendrites enable multi-task learning in dynamic environments. *Frontiers in Neuroinformatics*, 16:846219.
- Kaski, S., Kohonen, T. (1994). Winner-take-all networks for physiological models of competitive learning. *Neural Networks*, 7:973–984.
- Klasson, M., Kjellström, H., Zhang, C. (2023). Learn the time to learn: Replay scheduling in continual learning. *TMLR*.
- Kolen, J.F., Pollack, J.B. (1994). Backpropagation without weight transport. In *IEEE International Conference on Neural Networks (ICNN)*, pp. 1375–1380.
- Kubo, Y., Delanois, J.E., Bazhenov, M. (2025). Toward lifelong learning in equilibrium propagation: Sleep-like and awake rehearsal for enhanced stability. *arXiv:2508.14081*.
- Lin, E., Luo, W., Jia, W., Yang, S., Xu, W. (2026). Online continual learning via spiking neural networks with sleep enhanced latent replay. In *14th International Conference on Intelligent Control and Information Processing (ICICIP)*, pp. 360–367.
- Lässig, F., Aceituno, P.V., Sorbaro, M., Grewe, B.F. (2023). Bio-inspired, task-free continual learning through activity regularization. *Biological Cybernetics*, 117:345–361.
- Maeda, M., Miyajima, H. (1999). Competitive learning methods with refractory and creative approaches. *IEICE Transactions on Fundamentals*, E82-A(9):1825–1833.
- Mallya, A., Lazebnik, S. (2018). PackNet: Adding multiple tasks to a single network by iterative pruning. In *CVPR*.

373 Masse, N.Y., Grant, G.D., Freedman, D.J. (2018). Alleviating catastrophic forgetting using context-dependent  
374 gating and synaptic stabilization. *PNAS*, 115:E10467–E10475.

375 McKee, K., Hazy, T., Zheng, Y., Bugaud, Z., Miconi, T. (2026). Cortex-inspired continual learning: Unsupervised  
376 instantiation and recovery of functional task networks. *arXiv:2604.24637*.

377 Payeur, A., Guerguiev, J., Zenke, F., Richards, B.A., Naud, R. (2021). Burst-dependent synaptic plasticity can  
378 coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24:1010–1019.

379 Pham, Q., Liu, C., Hoi, S. (2021). DualNet: Continual learning, fast and slow. In *NeurIPS* 34.

380 Robinson, B.S., Lau, C.W., New, A., Nichols, S.M., Johnson, E.C., Wolmetz, M., Coon, W.G. (2022).  
381 Continual learning benefits from multiple sleep mechanisms: NREM, REM, and synaptic downscaling.  
382 *arXiv:2209.05245*.

383 Ross, G.J., Adams, N.M., Tasoulis, D.K., Hand, D.J. (2012). Exponentially weighted moving average charts for  
384 detecting concept drift. *Pattern Recognition Letters*, 33:191–198.

385 Rueckauer, B., Lungu, I.-A., Hu, Y., Pfeiffer, M., Liu, S.-C. (2017). Conversion of continuous-valued deep  
386 networks to efficient event-driven networks for image classification. *Frontiers in Neuroscience*, 11:682.

387 Saha, G., Garg, I., Roy, K. (2021). Gradient projection memory for continual learning. In *ICLR*.

388 Saper, C.B., Scammell, T.E., Lu, J. (2005). Hypothalamic regulation of sleep and circadian rhythms. *Nature*,  
389 437:1257–1263.

390 Shen, J., Ni, W., Xu, Q., Tang, H. (2024). Efficient spiking neural networks with sparse selective activation for  
391 continual learning. In *AAAI* 38, pp. 611–619.

392 Shen, J., Zhao, L., Xu, Q., Yang, Y., Chen, L., Pan, G., Chen, B. (2026). Robust selective activation with  
393 randomized temporal  $k$ -winner-take-all in spiking neural networks for continual learning. In *ICLR*.

394 Sorrenti, A., Bellitto, G., Proietto Salanitri, F., Pennisi, M., Palazzo, S., Spampinato, C. (2025). Wake–sleep  
395 consolidated learning. *IEEE Transactions on Neural Networks and Learning Systems*, 36(7):12668–12679.

396 Tadros, T., Krishnan, G.P., Ramyaa, R., Bazhenov, M. (2022). Sleep-like unsupervised replay reduces catastrophic  
397 forgetting in artificial neural networks. *Nature Communications*, 13:7742.

398 Tilley, M.J., Miller, M., Freedman, D.J. (2023). Artificial neuronal ensembles with learned context dependent  
399 gating. In *ICLR*.

400 Tononi, G., Cirelli, C. (2014). Sleep and the price of plasticity: From synaptic and cellular homeostasis to  
401 memory consolidation and integration. *Neuron*, 81:12–34.

402 Vitter, J.S. (1985). Random sampling with a reservoir. *ACM Transactions on Mathematical Software*, 11:37–57.

403 Vyazovskiy, V.V., Olcese, U., Hanlon, E.C., Nir, Y., Cirelli, C., Tononi, G. (2011). Local sleep in awake rats.  
404 *Nature*, 472:443–447.

405 Wang, S., Li, X., Sun, J., Xu, Z. (2021). Training networks in null space of feature covariance for continual  
406 learning. In *CVPR*.

407 Whittington, J.C.R., Bogacz, R. (2017). An approximation of the error backpropagation algorithm in a predictive  
408 coding network with local Hebbian synaptic plasticity. *Neural Computation*, 29:1229–1262.

## 409 A Proofs

410 *Proof of Proposition 1.* By induction on  $\ell$ , for every sample  $b$ . Assume  $a_{i,b}^{\ell-1}$  is unchanged (true  
 411 for  $\ell = 1$  since  $a_{i,b}^0 = x_b$ ). For any unit  $i$ ,  $\hat{z}_{i,b}^\ell - z_{i,b}^\ell = \sum_j \Delta_{ij}^\ell a_{j,b}^{\ell-1}$ , and each term vanishes: if  
 412  $j \in \mathcal{A}^{\ell-1}(X)$  then  $\Delta_{ij}^\ell = 0$ , otherwise  $a_{j,b}^{\ell-1} = 0$  for every  $b$ . Hence  $z_{i,b}^\ell$ , and  $a_{i,b}^\ell$  through (1), are  
 413 unchanged; the induction closes at the last hidden layer, and at the output layer whenever the readout  
 414 row satisfies the same condition.  $\square$

415 *Proof of Proposition 2.* A unit outside  $\mathcal{A}^\ell(X)$  contributes  $a_{i,b}^\ell = 0$  on every sample. After the update  
 416 its pre-activation on sample  $b$  is  $z_{i,b}^\ell + \delta_{i,b}$ ; under the stated margin it is still excluded by  $\Phi_k$  (or  
 417 rectified to zero), so  $a_{i,b}^\ell = 0$  still. The sample’s original winners keep their pre-activations, because  
 418 their synapses from active presynaptic units lie outside  $\Delta^\ell$  by the mask and their silent-presynaptic  
 419 terms vanish as in Proposition 1; so the same  $k$  units win with the same values, every other unit’s  
 420 input is untouched, and the layer’s activities are identical. When fewer than  $k$  units are positive on  
 421 sample  $b$ ,  $\tau_{k,b}^\ell = 0$  and the condition reads  $\sigma(z_{i,b}^\ell + \delta_{i,b}) = 0$ : a zero-valued entry carries no activity  
 422 whether or not  $\Phi_k$  selects it. The readout row is excluded from both propositions in the default system  
 423 because it is deliberately kept plastic.  $\square$

## 424 B One training step, the three channels and the skip path

425 Algorithm 1 lists one waking step of the default system with the state each quantity is read from.  
 426 Two conventions matter for the guarantees. First, the awake sets are read from the waking forward  
 427 pass, i.e. from the weights before the unmasked waking update, while the masked replay step acts  
 428 on the updated weights; Propositions 1–2 hold verbatim for a mask computed on the weights being  
 429 updated, and the drift reported in Section 3 is measured under the implemented order. Second,  
 430 the replay micro-batch is inferred from the suppression removed, so a refractory unit may fire on  
 431 replay and receive the update (Corollary 1 makes that update invisible to the waking batch for any  
 432 magnitude), and a burst of  $n_b$  micro-batches reuses the mask of its waking batch. The waking step is  
 433  $\eta_w = \eta m/256$  for a waking batch of  $m$  samples (the per-epoch drift of a batch-256 protocol), the  
 434 replay step  $\eta_r = 3\eta$  per micro-batch, with  $\eta = 0.02$  throughout.

---

**Algorithm 1** One waking step of the default system (rotation always on, one burst per batch).

---

**Require:** weights  $W_t$ , velocities  $v_t$ , suppression  $s(t)$ , buffer  $\mathcal{M}$ , pressures  $S$

- 1:  $a \leftarrow \text{forward}(X_t; W_t, s(t))$ ,  $\varepsilon \leftarrow \text{sweep}(a, y_t)$  Eqs. (1)–(2)
  - 2:  $(W, v) \leftarrow \text{unmasked heavy-ball step on } (W_t, v_t) \text{ with } G^\ell = \varepsilon^\ell (a^{\ell-1})^\top$ , rate  $\eta_w$  wake update
  - 3:  $\mathcal{A}^\ell \leftarrow \{i : \exists b, a_{i,b}^\ell \neq 0\}$  for every hidden layer;  $S_i \leftarrow S_i + \bar{a}_i(X_t)$  from the pre-update  $a$
  - 4:  $M^\ell \leftarrow \text{Eq. (3)}$ , unit mask  $\mathbf{1}[i \notin \mathcal{A}^\ell]$ , readout row and bias unmasked
  - 5:  $s_i^\ell(t+1) \leftarrow \mathbf{1}[i \in \mathcal{A}^\ell]$  refractory rotation
  - 6: write  $(X_t, y_t)$  into  $\mathcal{M}$  (reservoir)
  - 7: **for** each of the  $n_b$  replay micro-batches  $(X_r, y_r) \sim \mathcal{M}$  **do**
  - 8:  $a_r \leftarrow \text{forward}(X_r; W, s=0)$ ,  $\varepsilon_r \leftarrow \text{sweep}(a_r, y_r)$  suppression removed
  - 9:  $(W, v) \leftarrow \text{masked step, Eq. (4), with } G^\ell = \varepsilon_r^\ell (a_r^{\ell-1})^\top$ , rate  $\eta_r$ , masks  $M^\ell$
  - 10: **end for**
  - 11:  $S_i \leftarrow 0$  for every  $i \notin \mathcal{A}^\ell$  if a burst ran discharge
  - 12: **return**  $W_{t+1} = W$ ,  $v_{t+1} = v$ ,  $s(t+1)$
- 

435 **The three channels.** Table 2 lists the three ways a synapse enters the mask (3) and what each  
 436 guarantees; the rotation of Section 3 exists to move synapses from the second row into the third.

437 **A numerical example.** On one sample with  $k = 2$  and  $\sigma(z) = (3, 2, 0.5, 0)$ , units 3 and 4 are  
 438 asleep and  $\tau_k = 2$ . Any replay update with  $\sigma(z_3 + \delta_3) < 2$  and  $\sigma(z_4 + \delta_4) < 2$  leaves the activity  
 439 vector  $(3, 2, 0, 0)$  unchanged, because the two winners’ own inputs are not touched; if unit 4 is  
 440 refractory ( $s_4 = 1$ ) its bound is void and  $\delta_4$  may be arbitrarily large; and every synapse leaving unit 4  
 441 into the next layer may move by any amount, since its presynaptic activity is zero.

| Channel                                   | Condition                                        | Invisible to $X$                     | Source  |
|-------------------------------------------|--------------------------------------------------|--------------------------------------|---------|
| presynaptic silent                        | $j \notin \mathcal{A}^{\ell-1}(X)$               | for any magnitude                    | Prop. 1 |
| postsynaptic naturally asleep, active pre | $i \notin \mathcal{A}^{\ell}(X), s_i^{\ell} = 0$ | under the margin $\tau_{k,b}^{\ell}$ | Prop. 2 |
| postsynaptic suppressed                   | $s_i^{\ell} = 1$                                 | for any magnitude                    | Cor. 1  |

Table 2: **The three consolidation channels** opened by the synapse mask.

**Skip synapses.** For  $\ell \geq 3$  each deep hidden layer receives a bypass,  $z^{\ell} = W^{\ell}a^{\ell-1} + S^{\ell}a^{\ell-2} + b^{\ell}$ , with  $S^{\ell}$  under the same Dale projection and a feedback twin  $B_S^{\ell} \in \mathbb{R}^{n_{\ell+2} \times n_{\ell}}$  that carries the bypass target’s burst back to its source,

$$\varepsilon^{\ell} \leftarrow \varepsilon^{\ell} + \sigma'(z^{\ell}) \odot \left( (B_S^{\ell})^{\top} [\kappa \tanh(\varepsilon^{\ell+2}/\kappa) \odot \mathbf{1}(a^{\ell+2} > 0)] \right),$$

$S^{\ell}$  and  $B_S^{\ell-2}$  learned by the local product of the main text with the target layer’s decay  $\lambda_{\ell}$ , and replay-masked by the same rule two layers apart,  $M_{S,ij}^{\ell} = \mathbf{1}[j \notin \mathcal{A}^{\ell-2}(X) \vee i \notin \mathcal{A}^{\ell}(X)]$ , so that Propositions 1–2 extend unchanged.

## C The price of the substrate

Table 3 traces the path from a dense backpropagation MLP to the substrate used throughout, on i.i.d. MNIST. The cumulative ladder prices the entire constraint set at 1.2 points below the dense reference for  $4\times$  fewer synaptic operations per pass; the local burst-error objective itself costs nothing. Settling cannot merely be shortened ( $T=1$  collapses); the single-phase burst sweep matches five-step settling in one pass. The Kolen–Pollack decay and even restored weight transport are accuracy-neutral at this scale; the decay is useful as a regulariser on harder inputs. Spiking-flavoured training fails outright: spiking is affordable at inference, not as a training-time code. On the plain substrate SGD loses 1.3 points to Adam, while inside the consolidation loop the two tie (Table 1): per-synapse adaptive state is an asset while every synapse trains on every batch and no longer one once updates become mask-confined.

Table 3: **Substrate ablation** on i.i.d. MNIST (256–128 hidden core, three seeds; development phase, official test set). “Syn. ops” is the fraction of the dense MLP’s synaptic operations per forward pass. Upper block cumulative; lower block replaces one ingredient at its stage.

| Substrate                                                                  | Accuracy                      | Syn. ops  |
|----------------------------------------------------------------------------|-------------------------------|-----------|
| dense backpropagation MLP (autograd reference)                             | $97.7 \pm 0.1$                | 1.00      |
| local burst-error learner, unconstrained                                   | $97.8 \pm 0.1$                | 0.85      |
| + bounded error, Dale’s law, sign-concordant feedback (no transport)       | $97.1 \pm 0.0$                | 0.85      |
| + $k$ -WTA 10%                                                             | $97.0 \pm 0.2$                | 0.80      |
| + relaxation $T=5$                                                         | $97.0 \pm 0.1$                | 0.80      |
| + 30% distance-dependent connectivity                                      | $96.5 \pm 0.2$                | 0.24      |
| + burst gate and burst baseline                                            | $96.3 \pm 0.2$                | 0.24      |
| + single-phase burst sweep replacing relaxation ( <i>final substrate</i> ) | $96.5 \pm 0.1$                | 0.24      |
| random wiring instead of distance-dependent, equal density                 | $95.4 \pm 0.0$                | 0.24      |
| one settling step ( $T=1$ ) instead of five                                | $17.9 \pm 6.8$                | 0.80      |
| weight transport restored (mirroring, no KP decay)                         | $96.3 \pm 0.2$                | 0.24      |
| Kolen–Pollack decay removed                                                | $96.3 \pm 0.1$                | 0.24      |
| heavy-ball SGD instead of Adam ( $\eta = 0.02$ )                           | $95.2 \pm 0.3$                | 0.24      |
| $k$ -WTA 5% / 2% instead of 10%                                            | $95.9 \pm 0.1 / 94.2 \pm 0.2$ | 0.79/0.78 |
| multiplicative burst coding                                                | $78.8 \pm 1.9$                | 0.24      |
| training on $T_s=8$ spike counts                                           | $80.0 \pm 0.8$                | 0.24      |

## D Protocol details, cost accounting and the operating frontier

The V1-like front-end consists of 256 normalised random  $6\times 6\times 3$  patches sampled from the training images, applied with stride 2, rectified against each filter’s mean response over training images, and

quadrant-average-pooled (1024-D); nothing in it is learned. Split CIFAR-10 uses  $32 \times 32$  luminance so that distance-dependent wiring keeps its image geometry. The buffer uses reservoir writing; the night reference uses a 256–128 core, its preferred width in our sweeps. Runs are deterministic at a fixed CPU thread count.




Figure 4: **Training data flow.** Each waking batch is inferred under the current suppression mask; its awake sets define the synapse and unit masks for the replay micro-batch applied in the same step and the rotation mask for the next batch. The masked step confines the velocity and the decay to the mask; the readout row is the one deliberately plastic exception. Unit-level pressure decides when replay bursts run.

465

**Energy.** Converted to spiking inference ( $T_s=8$ ; thresholds calibrated on training images, evaluated on the held-out tenth), the default system runs at 92.7% (rate accuracy 94.2%) for a proxy of 117 nJ per sample against 2461 nJ for dense rate inference ( $21\times$ ) and 526 nJ for event-driven rate inference ( $4.5\times$ ), and is monotone in  $T_s$  (94.0% at  $T_s=32$ , 476 nJ); the rotation-free control reaches 94.2% at  $T_s=8$  and dips to 93.9% at  $T_s=32$  (Figure 6).

**Updates and wall-clock.** Samples are not the only cost. The headline schedule makes 16,885 masked updates against the night’s 480 ( $35\times$ ; counts from the reported runs, whose stream is nine tenths of the training set), and its training wall-clock on one CPU thread is 8.8 min against 1.4 for the night, 1.0 with no replay and 8.6 for unmasked interleaved replay (one run each, same protocol and thread count). The per-batch replay step, not the sample count, is the online price of consolidating during the stream, in an unoptimised implementation whose mask construction is dense. Replay cost is reported in samples ( $R = \text{replay batches} \times \text{batch size}$ ). The night costs  $480 \times 256 \approx 1.2 \times 10^5$  samples per run. The headline configuration (one 16-sample micro-batch per waking batch) costs  $16,885 \times 16 \approx 2.7 \times 10^5$  ( $2.2\times$  the night) for  $91.6 \pm 0.3$ . Two schedules meet the night’s budget within  $1.1\times$ : eight-sample micro-batches at every waking batch ( $1.35 \times 10^5$ ,  $91.1 \pm 0.4$ ) and 16-sample micro-batches at every second one (8,442 updates,  $1.35 \times 10^5$ ,  $91.2 \pm 0.2$  sequential,  $93.1 \pm 0.2$  static); pressure-triggered bursts of 16 at a third of the events reach  $88.6 \pm 1.6$ . Accuracy therefore depends on the *frequency* of small isolated updates rather than on the number of replayed samples. Figure 5 draws the operating frontier of the default substrate.

**Baseline tuning and paired comparison.** The offline-night reference was swept over replay batch count (20, 40 per night), replay batch size (16, 64, 256), learning-rate gain ( $3\times$ ), core width (256–128 and 512–256), buffer size (200, 1000, 5000) and timing (clocked after each epoch, novelty-timed, surprise-timed), and its best sequential cell is reported; BP+ER was swept over learning rate ( $3 \cdot 10^{-4}$ ,  $10^{-3}$ ,  $3 \cdot 10^{-3}$ ), width and buffer size; the local learner over  $\eta$  (Appendix H), replay batch size, cadence and the gates of Table 1. Selection everywhere was sequential accuracy at  $K=1000$  with the static axis reported alongside, on the official test split (the development set); the comparisons reported are on the held-out split (Appendix L). Pairing seeds by index on that split, the local system’s margin over the narrow-substrate night is +2.5 points (bootstrap 95% CI [+1.3, +4.8], three seed pairs), over the same-substrate night +2.6 ([+0.2, +5.2], six pairs; the night’s seed variance makes this one fragile, and on the second split it shrinks to +0.3), over unmasked interleaved ER +6.3






Figure 5: **Operating frontier** at the 10% activity fraction under heavy-ball SGD: split-MNIST sequential accuracy against i.i.d. static accuracy (held-out split). Blue: non-dominated by seed means (the bout-committed gate, the silent mask, unmasked rotation-free replay); purple: dominated. The top cluster (always-on rotation, bout gate, anchor, masked Adam) lies within seed noise of one another; the frontier trades sequential for static accuracy through the silent and unmasked variants, and the adaptive decay is dominated on this split.

496  $([+3.8, +9.2])$ , over BP+ER +2.8  $([+2.5, +3.2])$ , over the silent mask +3.6  $([+3.1, +4.1])$  and over  
 497 rotation alone +0.1  $([-0.3, +0.5])$ .




Figure 6: **Left:** the substrate decomposition on feature CIFAR-10, transplanting the local network’s ingredients into BP+ER one at a time (each cell at its better of two learning rates; three seeds). **Right:** spiking inference, i.i.d. MNIST accuracy against the idealised 45-nm FP32 arithmetic-energy proxy for  $T_s \in \{4, 8, 16, 32\}$ , with the dense (square) and event-driven (diamond) rate references, held-out split; the default is monotone in  $T_s$ , the rotation-free control dips 0.4 at  $T_s=32$ .

## 498 E Rotation gates at the 5% activity fraction

499 Appendices E–H report development-phase numbers (official test set) from the earlier 5% phases and  
 500 the learning-rate sweeps; they were not re-run under the held-out protocol and support mechanism  
 501 orderings only. At the 5% fraction of the earlier phases, always-on rotation cost 1.3–3.0 points on  
 502 i.i.d. MNIST (95.0  $\rightarrow$  92.0–93.7). Gating rotation by unit-level pressure only trades the axes (static  
 503 91.6  $\rightarrow$  94.0 as sequential 88.8  $\rightarrow$  80.7). The relative-novelty gate removes the price at no sequential  
 504 cost on MNIST (static  $95.0 \pm 0.4$  with rotation on 1.4–17% of batches; sequential 90.3–90.9). On  
 505 split CIFAR-10 the price reverses, with rotation gaining +2.5 to +6.5 static points in the underfit  
 506 regime, while the pure novelty gate misfires there (sequential  $19.1 \pm 1.1$ : fast/slow  $\rightarrow$  1 on a stream  
 507 that never becomes predictable). The unmastered clause resolves this with one setting across all four  
 508 regimes ( $27.2 \pm 0.7$  /  $25.3 \pm 1.0$  on CIFAR,  $90.4 \pm 0.9$  /  $94.3 \pm 0.2$  on MNIST). Those gate numbers  
 509 use masked Adam; under heavy-ball SGD the per-batch gate’s flicker destroys the optimiser’s gain,

and the gate must be committed to minimum-duration bouts ( $M=2048$  batches, loosely modelled on the consolidated bouts of the sleep–wake switch [Saper et al., 2005]; Table 1). Figure 7 shows the three regimes.




Figure 7: **Rotation at 5%: price, trade-off and gate.** Left: on i.i.d. MNIST the novelty gate restores no-rotation accuracy. Middle: on split-MNIST rotation is required and the gate keeps it. Right: on split CIFAR-10 rotation is a gain on both axes.

## F Probes of the residual static price

**Bout length and the OFF-side refractory.** Bout length is a genuine control parameter:  $M=512$  gives  $91.7 \pm 0.4/94.0 \pm 0.5$ , recovering most of what the per-batch gate loses, while  $M=4096$  drives the static duty to 0.94 and the refund vanishes. Completing the switch with an OFF-side refractory (triggers ignored for  $M_{\text{off}}$  batches after a bout) is a clean null result: short refractories never engage on the static stream, and a long one (4096) lowers the static duty to 0.53 without moving static accuracy ( $94.0 \pm 0.7$ ) while monotonically eroding the sequential axis ( $92.1 \rightarrow 91.0$ ). The residual static price sits in the rotation of settled coalitions itself, not in how often the gate fires.

**Structural exemption and synaptic anchoring.** Sparing the top- $q$  units by long-term use from ever being benched traces a smooth trade-off ( $q=0.10$ :  $90.9 \pm 0.1/94.2 \pm 0.2$ ;  $q=0.25$ :  $89.1 \pm 0.4/95.4 \pm 0.2$ ) that the bout gate dominates: temporal commitment beats structural exemption. Two-timescale synapses in the spirit of synaptic consolidation cascades [Benna and Fusi, 2016], with a slow anchor absorbing the fast weight at rate  $\mu$  while the fast weight decays toward it at rate  $\lambda$ , ticking only on waking updates, do not dissolve the price (no cell exceeds the anchor-free static accuracy); weak symmetric coupling acts as a sequential stabiliser at 5% ( $92.4 \pm 0.2$ , variance halved) and is neutral at 10% ( $92.7 \pm 0.4/94.6 \pm 0.3$ ).

## G Small buffers and sleep-anchored down-selection

At  $K=200$  the night at first appeared to keep a 2.6-point edge; on the same 512–256/5% substrate the night reaches  $84.0 \pm 3.4$  while noisy pressure-burst local sleep reaches  $84.3 \pm 1.0$  and local+night  $84.8 \pm 0.9$ , i.e. parity, with three-fold smaller variance for the wake-side scheme. The strongest tiny-buffer number remains the narrower night (86.9), so the offline window retains an advantage only in combination with narrow substrates. Replay diversity (input noise  $\sigma=0.2$ ) is a wake-side lever: it lifts every local scheme ( $81.8 \pm 2.9 \rightarrow 84.3 \pm 1.0$ ) and hurts the night. Unmasked, rotation-free waking replay collapses ( $71.0 \pm 1.7$ ). Anchoring structural plasticity (prune the weakest live synapses, regrow with a distance bias) to consolidation windows, motivated by synaptic homeostasis [Tononi and Cirelli, 2014], is free on the continuous substrate ( $90.9 \pm 0.2$  vs.  $90.8 \pm 0.1$ ), rescues the wide-sparse night ( $87.2 \pm 3.8$  vs.  $84.1 \pm 1.3$ ), and harms the episodic-burst substrate ( $84.6 \pm 0.5$  vs.  $89.9 \pm 0.1$ ): down-selection, like replay, needs either density or an offline window.

## H Learning-rate sensitivity

At the 5% fraction, under always-on rotation, SGD holds a sequential plateau over  $\eta \in [0.01, 0.02]$  ( $91.9 \pm 0.5/92.0 \pm 0.6$ ) and a static plateau over  $[0.02, 0.05]$  ( $93.3\text{--}93.7$ ), decaying gracefully outside ( $86.3/89.2$  at  $\eta=0.1$ ); masked Adam, swept over a tenfold range, has its sequential optimum at the default ( $90.8 \pm 0.1$  at  $10^{-3}$ ;  $88.6$  at  $3 \cdot 10^{-4}$ ,  $88.9$  at  $3 \cdot 10^{-3}$ ), and no Adam setting reaches the heavy-ball joint point. Unmasked SGD is poor and unstable at every rate ( $84\text{--}88$ , seed deviations  $1.8\text{--}4.1$ ). The same two-sided sweep at the 10% fraction (three seeds) gives the same result: heavy-ball SGD holds  $92.2 \pm 0.7/93.3 \pm 0.2$  at  $\eta=0.01$ ,  $92.7 \pm 0.3/94.7 \pm 0.3$  at  $0.02$  (the joint optimum) and  $91.5 \pm 0.6/94.7 \pm 0.3$  at  $0.05$ , decaying at  $0.005$  ( $91.3/91.3$ ) and destabilising at  $0.1$  ( $79.5 \pm 17.3$ ); masked Adam has its optimum at the default ( $91.8 \pm 0.6/94.6 \pm 0.2$  at  $10^{-3}$ ;  $89.9/89.9$  at  $3 \cdot 10^{-4}$ ,  $90.9/94.3$  at  $3 \cdot 10^{-3}$ ). The optimiser comparison of Table 1 is made at each optimiser’s best rate.

## I Momentum, exactness and other negative results

**Momentum.** With momentum zero the default replay step ( $3\eta$  per micro-batch against the batch-scaled waking step  $\eta m/256$ , Appendix B) diverges or starves hidden activity at every learning rate tested ( $\eta \in [0.02, 0.4]$ :  $\approx 10\%$ , hidden dimensionality 0); the waking path alone learns normally at the matched effective step, so the collapse is the replay path’s: the velocity’s  $\sim 10\times$  averaging of the 16-sample replay micro-batches is what heavy-ball provides. Re-scaling the replay gain ( $\eta=0.2$ , gain 0.3) makes pure SGD viable at  $88.4 \pm 0.5/89.4 \pm 0.3$ ; other settings are fragile (development-phase sweep).

**Exactness measurement.** The numbers in Section 3 are population rates over every replay step of a full run under the reported protocol; the rest of this paragraph reports development runs (official test split, full training set, 18,760 updates), on which the default configuration’s hidden-code rate was 0.32%. Under the 5% configuration with masked Adam, the isolated update altered the top hidden code of 1.0% of waking samples per step; on raw CIFAR-10 (one seed) the hidden-code rate is 0.46% while the prediction rate rises to 3.6%, because low-margin predictions flip more easily through the plastic readout. The rates are taken under the implemented order of Algorithm 1, in which the awake sets come from the activities before the waking update. Re-inferring the batch after that update and masking on the post-update awake sets lowers the hidden-code rate from 0.32% to 0.30% (0.17% with the readout isolated) and leaves the prediction rate (0.30%; 0.000% isolated), the margin-violation rate and the accuracy unchanged ( $92.2$  vs.  $92.0$ , one seed), so the residual drift comes from the unenforced margin, not from the ordering.

**Replay gain.** Sweeping the replay gain over  $\{1/16, 1/4, 1, 3, 10\}$  (default 3;  $\eta_r = 3\eta$  against a waking step of  $\eta/16$ ) moves masked replay  $67.2 \pm 3.0 \rightarrow 86.8 \pm 1.1 \rightarrow 90.8 \pm 0.1 \rightarrow 91.6 \pm 0.3 \rightarrow 74.2 \pm 32.1$  and unmasked replay without rotation  $61.6 \pm 7.2 \rightarrow 81.4 \pm 4.5 \rightarrow 85.7 \pm 2.2 \rightarrow 85.4 \pm 3.5 \rightarrow 47.3 \pm 30.7$  (sequential; three seeds, six at gain 10): masked replay peaks at the default, unmasked replay plateaus from gain 1 to 3, gain 10 destabilises both (two of six seeds collapse for the masked system, four for the unmasked), and no replay gain makes the unmasked control competitive.

**Single-pass streams.** With one epoch per task instead of five (everything else as in the headline cells, three seeds), refractory local sleep reaches  $91.8 \pm 0.2$  ( $F=5.7$ ) against  $91.6$  with five epochs, BP+ER  $88.7 \pm 0.3$ , the offline night  $75.1 \pm 3.5$  (one night per task is too few), unmasked interleaved replay  $76.0 \pm 20.1$  and no buffer 18.1: the proposed system keeps its lead over every tested control in a single pass, while the offline night falls below BP+ER.

**Rejected under matched budgets.** The following alternatives were tested and found not to help: a frozen dentate-gyrus-style sparse expansion front-end (input decorrelation does not propagate through learned  $k$ -WTA layers: input task overlap drops  $0.79 \rightarrow 0.18$ , yet every downstream metric worsens by 3–11 points); generative (REM) replay as a substitute for episodic samples; margin-based reverse learning; surprise-gated memory writing; multiplicative burst coding; absolute-threshold rotation gates; soft (gain-reduction) rotation; mirror-direction isolation; OFF-side gate refractoriness; strongly coupled synaptic anchors; utility-exempt rotation.

## 590 J The controller: composition rules, CIFAR-100 and reproducibility

Table 4: **One dataset-independent ratio target replaces per-dataset decay tuning.** Sequential / static accuracy under the tuned fixed decay ( $10^{-3}$ ;  $10^{-4}$  with  $3\times$  targets as the CIFAR-100 rescue) and the controller (5) at  $\rho = 0.25$  everywhere (drive form). Full system in every cell, held-out split, three seeds (six on split-MNIST); CIFAR-100 numbers are floor-regime and reported for ordering only (BP+ER 6.8 raw, 10.1 with features).

|                  | fixed $\lambda$ (tuned) |                | controller     |                |
|------------------|-------------------------|----------------|----------------|----------------|
|                  | seq.                    | static         | seq.           | static         |
| split-MNIST      | $91.6 \pm 0.3$          | $94.1 \pm 0.2$ | $90.6 \pm 0.6$ | $93.5 \pm 4.3$ |
| split CIFAR-10   | $28.9 \pm 1.0$          | $28.0 \pm 1.2$ | $26.2 \pm 1.1$ | $30.1 \pm 0.2$ |
| feature CIFAR-10 | $42.3 \pm 0.6$          | $42.4 \pm 1.1$ | $39.1 \pm 0.4$ | $43.6 \pm 0.6$ |
| split CIFAR-100  | $1.2 \pm 0.1$           | $1.0 \pm 0.1$  | $3.2 \pm 0.8$  | $4.1 \pm 0.6$  |
| + features       | $1.1 \pm 0.3$           | $1.1 \pm 0.2$  | $5.2 \pm 1.6$  | $6.4 \pm 0.6$  |

591 A gradient-referenced form of the controller ( $\rho \eta \text{EMA}[\overline{|g^\ell|}]/\overline{|W^\ell|}$ ) behaved equivalently in devel-  
 592 opment (official test set, not re-run: split-MNIST  $92.2 \pm 0.4/95.2 \pm 0.3$  against  $91.7/96.1$  for the  
 593 drive form there). The drive estimate must see only unamplified waking updates: inflated one-hot  
 594 targets and the offline night’s  $3\times$ -gain batches push the EMA to its clip and erase day learning  
 595 (development phase: night  $\times$  controller  $60.2 \pm 2.8$  against  $90.9 \pm 3.2$  for the night alone), whereas  
 596 anchor and bout gate compose cleanly. The controller does not subsume  $\lambda$ ’s role as a capacity  
 597 regulariser against overfitting on small tabular problems, information no drive ratio contains. On  
 598 CIFAR-100 (ten tasks of ten classes; every method at floor, chance 1%, BP+ER  $10.1 \pm 0.4$  with  
 599 features and  $6.8 \pm 0.2$  raw, with twice the local learner’s forgetting) depth is a net loss that skips only  
 600 half refund (development-phase numbers: four hidden layers  $1.9/1.6 \rightarrow 3.8/3.0$  with skips against  
 601  $5.7/6.0$  at two), while width at the short chain length lifts the local system to  $5.8 \pm 0.9/7.7 \pm 1.1$   
 602 ( $1024\text{--}512$ ). The thin-signal regime wants capacity, not depth, and the static-axis wall (6–8% under  
 603 every substrate change) indicates that the front-end is the binding constraint there. All runs reproduce  
 604 bitwise at a fixed CPU thread count; a different thread count changes the reduction order, which  
 605  $k$ -WTA’s discontinuity amplifies into a different trajectory.

## 606 K Isolation across replay batch sizes

607 Figure 8 extends the left panel of Figure 2 to two-sample replay micro-batches and shows, per batch  
 608 size, the seed-paired margin of the full system over rotation with unmasked replay (held-out split;  
 609 six seeds at batches 2, 4 and 16, three at 8). From 16 down to 4 the margin is  $0.1 [-0.3, +0.5]$ ,  $0.3$   
 610  $[-0.2, +0.6]$  and  $0.3 [-0.5, +1.2]$ : isolation costs nothing and buys nothing measurable in mean  
 611 accuracy while the rotation keeps the live coalition tolerant of unmasked replay. At batch 2 the replay  
 612 gradient is too noisy for either system. The isolated one falls to  $79.0 \pm 3.2$  (forgetting  $21.5 \pm 4.2$   
 613 against 6.4 at batch 16), with every seed between 75.5 and 84.3; the unmasked one bifurcates, with  
 614 three seeds finishing at 82.5–84.8, slightly *above* their isolated counterparts (+0.3, +1.8, +8.1), one  
 615 at 58.2 and two at chance (9.9). Leaking replay onto the live coalition is therefore not a pure loss,  
 616 since the surviving seeds show that the extra plasticity can help, but it exposes a severe-failure mode  
 617 (final accuracy below 60%: two seeds at chance, one at 58.2) that the isolated system did not show in  
 618 six seeds. With the current input’s hidden computation held invariant, the noisiest replay degraded the  
 619 system gracefully in every seed tested. Six seeds without a severe failure bound its frequency but do  
 620 not prove the mode absent, and the masked system itself destabilises at replay gain 10 (Appendix I).  
 621 The paired mean at batch 2 (+24  $[-0.2, +48.5]$ ) describes the unmasked failures, not a margin, and  
 622 is kept out of the ranges quoted in the main text.

## 623 L Development-phase numbers and a second held-out split

624 Configurations were selected on the official test splits (the development set) and frozen; the compar-  
 625 isons in the main text re-train them on nine tenths of the training data and evaluate on the held-out  
 626 tenth, which took part in no selection. Table 5 places the development-set numbers beside the  
 627 held-out ones and beside a second held-out tenth for the headline rows (drawn with a different seed;




Figure 8: **Isolation across replay batch sizes** (held-out split). **Left:** split-MNIST accuracy against replay batch size for isolated replay (the full system) and for rotation with unmasked replay, every seed shown; at batch 2 no unmasked mean is drawn because three of six seeds fail severely (values at the axis floor). **Right:** the seed-paired margin of isolation with a 95% bootstrap interval at the batch sizes where the unmasked run does not collapse.

the two tenths share 9% of their samples, so they are not disjoint test sets). Over the 87 sequential configurations run under both protocols, reported numbers sit 1.1 points below the development numbers on average and rank-correlate with them at Spearman  $\rho = 0.97$  (a robustness check, not a proof that selection bias is absent); every ordering carried by the paper holds on both held-out splits, with one magnitude that does not: the same-substrate night, bimodal on the first split ( $89.0 \pm 3.5$ ), reaches  $91.5 \pm 0.5$  on the second and trails the full system by 0.3 there. The adaptive decay’s static margin on MNIST is likewise split-dependent ( $93.5 \pm 4.3$  with one collapsed seed on the first split,  $95.2 \pm 0.4$  on the second). Table 6 gives the per-task accuracy matrix of the headline configuration on the reported split.

Table 5: **Development-set (official test split) numbers beside the held-out numbers and a second held-out split.** Sequential accuracy, with the static axis on its own row where the paper reports it; six seeds on the reported split for the headline rows, three elsewhere. Lower block: raw split CIFAR-10.

|                                                 | development set | held-out split 1 | held-out split 2 |
|-------------------------------------------------|-----------------|------------------|------------------|
| always-on rotation (default)                    | $92.7 \pm 0.3$  | $91.6 \pm 0.3$   | $91.8 \pm 0.4$   |
| static                                          | $94.7 \pm 0.3$  | $94.1 \pm 0.2$   | $94.1 \pm 0.6$   |
| bout-committed gate                             | $92.3 \pm 0.2$  | $91.8 \pm 0.3$   | $92.0 \pm 0.6$   |
| rotation alone (unmasked replay)                | $92.1 \pm 0.5$  | $91.5 \pm 0.3$   | $91.8 \pm 0.6$   |
| static                                          | $94.7 \pm 0.1$  | $93.7 \pm 0.3$   | $93.7 \pm 0.3$   |
| silent mask (no rotation)                       | $88.8 \pm 0.8$  | $88.0 \pm 0.7$   | $88.6 \pm 1.3$   |
| unmasked interleaved ER (SGD)                   | $87.0 \pm 1.1$  | $85.4 \pm 3.5$   | $80.0 \pm 14.1$  |
| BP + ER                                         | $89.3 \pm 0.8$  | $88.8 \pm 0.3$   | $88.5 \pm 0.7$   |
| offline night, same substrate                   | $90.9 \pm 3.2$  | $89.0 \pm 3.5$   | $91.5 \pm 0.5$   |
| offline night, narrow substrate                 | $90.7 \pm 0.4$  | $89.2 \pm 1.7$   | $89.3 \pm 0.6$   |
| adaptive $\lambda$ (drive)                      | $91.7 \pm 0.2$  | $90.6 \pm 0.6$   | $90.8 \pm 0.1$   |
| static                                          | $96.1 \pm 0.2$  | $93.5 \pm 4.3$   | $95.2 \pm 0.4$   |
| five hidden layers + skips + adaptive $\lambda$ | $91.5 \pm 0.4$  | $91.0 \pm 0.3$   | $91.1 \pm 0.3$   |
| static                                          | $96.2 \pm 0.3$  | $95.8 \pm 0.3$   | $95.7 \pm 0.1$   |
| no buffer                                       | $19.6 \pm 0.0$  | $19.4 \pm 0.0$   | $19.5 \pm 0.0$   |
| <i>split CIFAR-10 (raw)</i>                     |                 |                  |                  |
| rotation                                        | $29.5 \pm 0.8$  | $28.9 \pm 1.0$   | $29.7 \pm 0.9$   |
| offline night                                   | $26.2 \pm 0.6$  | $25.7 \pm 1.2$   | $26.5 \pm 0.4$   |
| BP + ER                                         | $25.1 \pm 1.5$  | $24.8 \pm 0.3$   | $25.8 \pm 0.8$   |
| unmasked ER                                     | $17.0 \pm 6.1$  | $15.4 \pm 4.2$   | $15.5 \pm 4.5$   |
| no buffer                                       | $16.4 \pm 0.2$  | $16.3 \pm 0.1$   | $16.6 \pm 0.0$   |

Table 6: **Per-task accuracy matrix** of the headline configuration on split-MNIST (held-out split, mean over six seeds): row = after training task  $i$ , column = accuracy on task  $j$ . Forgetting  $F$ , the mean of the diagonal minus the last row over tasks 1–4, is  $6.4 \pm 0.6$ ; the drop is concentrated on tasks 2–3 and the first task is retained.

| after task | T1   | T2   | T3   | T4   | T5   |
|------------|------|------|------|------|------|
| 1          | 99.8 |      |      |      |      |
| 2          | 98.7 | 97.2 |      |      |      |
| 3          | 97.7 | 92.7 | 97.0 |      |      |
| 4          | 97.5 | 89.6 | 92.9 | 97.5 |      |
| 5          | 96.7 | 88.1 | 89.7 | 91.5 | 91.8 |