
Inference as Consolidation: Continual Learning with No Offline Phase by Replay in the Silent Degrees of Freedom

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 An agent that keeps learning in deployment may not be able to afford dedicated
2 consolidation downtime, yet replay-based continual learning almost always relies
3 on offline rehearsal phases or on interleaving replayed samples with the incoming
4 stream. Brains face the same constraint and solve part of it with *local sleep*:
5 brief, use-dependent off-periods of individual circuits during wakefulness. We ask
6 whether a network trained by local, biologically constrained rules can consolidate
7 *during inference*, with no offline phase at all. We introduce (i) an *isolation* rule
8 that confines replay updates to hidden synapses invisible to the current input
9 under *k*-winner-take-all dynamics — exact for silent presynaptic and suppressed
10 units, near-exact elsewhere by measurement — (ii) a *refractory rotation* rule under
11 which units that have just fired sit out the next competition, widening the isolated
12 channel, and (iii) internal control signals: a unit-level homeostatic pressure that
13 triggers replay bursts and a relative-novelty gate. The construction inverts the usual
14 direction of non-interfering continual learning: the hidden computation is held
15 invariant — exactly on the proven channel, near-exactly end to end — while past
16 memories are consolidated through degrees of freedom the current batch leaves
17 unused. On class-incremental split-MNIST the system reaches $92.7 \pm 0.3\%$ with no
18 offline phase, above the best offline-rehearsal schedule ($90.7 \pm 0.4\%$), at $2.4\times$ its
19 replay samples; within $1.2\times$ of that budget, eight-sample micro-batches still reach
20 $91.8 \pm 0.3\%$, where unmasked replay without rotation and offline rehearsal degrade
21 sharply; the ordering transfers to split CIFAR-10 and replicates on a held-out split
22 never used for selection. A drive-referenced controller for the rule’s synaptic decay
23 replaces per-dataset tuning with one dataset-independent ratio target and, with
24 locally learned skip synapses, holds a five-hidden-layer stack at the two-layer level
25 where the tuned fixed decay without skips collapses.

1 Introduction

26
27 Foundation-model and embodied agents are increasingly expected to keep learning after deployment,
28 from streams that are non-stationary and never task-labelled. The dominant tool against forgetting is
29 replay from an episodic memory, and nearly every replay-based method consumes it in one of two
30 ways: in a dedicated offline phase — a “night” during which the agent stops acting and rehearses
31 — or by interleaving replayed samples with the live stream. Both cost an agent. An offline phase is
32 downtime during which the system is unavailable or silently changing; interleaved replay perturbs
33 the computation currently serving the input, and grows fragile as replay batches shrink toward what
34 an online budget allows.

Biology suggests a third option. Sleep-like off-periods occur *locally*, in individual cortical circuits of awake animals [Vyazovskiy et al., 2011], and inducing such ON/OFF alternation in one hemisphere of awake mice discharges sleep pressure and rescues memory consolidation, whereas an equal tonic reduction of firing does not [Driessen et al., 2026]. Some core functions of sleep can thus be discharged locally during wakefulness. We take this as an algorithmic motivation — not as a claim about where biological consolidation occurs — to ask whether the computational degrees of freedom a learner is not using can host consolidation while it acts, and what that costs.

We study the question in a deliberately constrained learner — a two-hidden-layer network trained end to end by local, biologically motivated rules (no weight transport, bounded burst-like error channel, Dale’s law, sparse population codes, sparse distance-dependent wiring, single forward–feedback sweep) — because those constraints expose exactly which degrees of freedom the current input does and does not use. The contributions are:

- **Isolated replay during inference.** A synapse may take a replay update only if its presynaptic unit is silent for the current input or its postsynaptic unit is asleep. Under k -WTA dynamics this leaves the hidden computation exactly invariant for silent-presynaptic and suppressed units (Propositions 1–2); we measure how nearly it holds elsewhere and end to end. With rotation on, replay micro-batches of 8–16 samples suffice, where unmasked replay without rotation and offline schedules degrade; isolation adds 0.6–0.7 points there and the guarantee, and replay confined to the readout is at chance: the effect is hidden-layer consolidation.
- **Refractory rotation.** Units that fired on the current batch sit out the next competition. This alternation — the algorithmic counterpart of the ON/OFF requirement in Driessen et al. [2026] — widens the consolidable channel from ≈ 0.25 to 0.53 of the replay-active units and is worth +3.9 sequential points on top of isolation (+5.1 alone); a random alternation of matched size recovers most of that, use-dependence the remaining 1.9; a tonic reduction, like the biological control, loses on both axes.
- **Internal triggers** (schedulers for sparse replay budgets; the headline system replays at every batch). Replay is triggered by a unit-level homeostatic pressure; the tested surprise trigger fails during wake, whereas offline rehearsal is best timed by novelty in our sweeps. A relative-novelty gate on rotation is scale-free where absolute thresholds fail across the regimes considered here.
- **One ratio target instead of per-dataset decay tuning, and depth.** A drive-referenced controller for the rule’s synaptic decay removes its last per-dataset constant and, with locally learned skip synapses, holds a five-hidden-layer stack at the two-layer level on both axes where the tuned fixed decay without skips collapses to chance.

Everything is evaluated on two axes at once — the class-incremental stream and the identical learner on the i.i.d. shuffle of the same data — because consolidation during inference must not damage ordinary learning. The scale is small (MLPs on MNIST/CIFAR variants); the object of study is the mechanism, whose assumptions (sparse codes, a local error signal) are the ones a larger system would have to satisfy to inherit the guarantee.

2 Related work

Sleep-inspired consolidation. The Bazhenov line implements sleep as a global offline phase driven by noise under local plasticity [Tadros et al., 2022, Bazhenov et al., 2026, Golden et al., 2022, Kubo et al., 2025]; wake–sleep frameworks with NREM/REM stages [Sorrenti et al., 2025, Robinson et al., 2022] and systems such as SIESTA [Harun et al., 2023] likewise schedule consolidation offline, and SESLR [Lin et al., 2026] appends a noise-enhanced sleep phase (a frozen-extractor spiking classifier training only on its buffer) to correct online learning’s recency bias. In all of these sleep is global, offline and clocked; none consolidates in currently unused units during inference, and the bias correction such a phase delivers is what our plastic readout receives continuously from isolated replay.

Gating and history-dependent competition. Context-dependent gating activates a subnetwork per task from task identity [Masse et al., 2018]; learned gates [Tilley et al., 2023], dendritic context [Iyer et al., 2022], task-free sparsification [Lässig et al., 2023] and recovered functional masks [McKee et al., 2026] derive the gate from the input. Refractory competition has classical precedents [Kaski and Kohonen, 1994, Maeda and Miyajima, 1999], and spiking continual learners raise firing thresholds

with use [Shen et al., 2024]. Our rotation differs in purpose: recently active units are benched for one competition to enlarge a *replay channel*, not to allocate resources across tasks.

Null-space and parameter isolation. A mature line protects *previous* tasks: gradient or activation-subspace projections [Farajtabar et al., 2020, Saha et al., 2021], feature-covariance null spaces [Wang et al., 2021], k -winner sparsity with such projections [Abbasi et al., 2022], frozen subnetworks [Golkar et al., 2019, Mallya and Lazebnik, 2018]. We invert the protected direction: the ongoing computation is held invariant while old-memory replay is written into degrees of freedom the current sparse support exposes as invisible — no stored bases or task masks, and the “null space” changes every batch for free. McKee et al. [2026] and Bricken et al. [2023] note that optimiser state can break such guarantees; ours is mask-confined. Complementary-learning-system methods [Arani et al., 2022, Pham et al., 2021] raise the value of each replayed sample but still replay at every step; learned replay scheduling [Klasson et al., 2023] is related to our trigger study.

3 Method

Setting. A stream of labelled examples is presented in T contiguous tasks with disjoint class sets; no task identity is available at test time; a single shared C -way readout is evaluated on all classes after the last task (class-incremental learning). We report final accuracy A_{seq} , forgetting F , the accuracy A_{iid} of the identical learner and budget on the i.i.d. shuffle (the *static* axis), the replay cost R in samples, and an arithmetic-energy proxy E ; the learner is confined to the cortical constraint set $\mathcal{B}$ above and to an episodic memory of K raw examples with reservoir writing [Vitter, 1985]. The question is whether the offline term of consolidation can be removed entirely without losing A_{seq} , and at what price in A_{iid} and R .

Base learner. Layer activities are $a^0 = x$ and, for $\ell = 1, \dots, L$,

$$z^\ell = W^\ell a^{\ell-1} + b^\ell, \quad a^\ell = \Phi_k(\sigma(z^\ell) \odot (1 - s^\ell)), \quad (1)$$

where σ is a rectifier, $s^\ell \in \{0, 1\}^{n_\ell}$ the suppression mask of the rotation rule below, and Φ_k keeps the k largest entries of each row (k -WTA); a per-unit gain exists in the implementation for homeostatic scaling but is fixed at 1 throughout. Layers $1, \dots, L-1$ are hidden; layer L is a linear C -way readout whose pre-activation z^L is the output, scored by squared error against the one-hot target y . Errors come from one top-down sweep through learned feedback weights $B^\ell \in \mathbb{R}^{n_{\ell+1} \times n_\ell}$ (no weight transport),

$$\varepsilon^L = y - z^L, \quad \varepsilon^\ell = \sigma'(z^\ell) \odot \left((B^\ell)^\top [\kappa \tanh(\varepsilon^{\ell+1}/\kappa) \odot \mathbf{1}(a^{\ell+1} > 0)] \right), \quad (2)$$

a bounded ($|\varepsilon_i^\ell| \leq \kappa \sum_j |B_{ji}^\ell|$ per component), event-gated burst signal in the spirit of predictive-coding and burst-dependent learners [Whittington and Bogacz, 2017, Payeur et al., 2021]. Forward and feedback weights follow the same local product with a shared decay λ (Kolen–Pollack alignment [Kolen and Pollack, 1994, Akrouit et al., 2019]), $\Delta W^\ell \propto \varepsilon^\ell (a^{\ell-1})^\top - \lambda W^\ell$, $\Delta B^{\ell-1} \propto \varepsilon^\ell (a^{\ell-1})^\top - \lambda B^{\ell-1}$, under Dale’s law with sign-concordant feedback. The optimiser is heavy-ball SGD — one per-synapse velocity, no second-moment state (Section 4 shows why neither Adam nor momentum-free SGD serves inside the consolidation loop). The base learner costs 1.2 points against a dense backpropagation MLP on i.i.d. MNIST at $4\times$ fewer synaptic operations (Appendix C).

Isolated replay. Let $X = \{x_b\}_{b=1}^m$ be the current waking batch, $a_{i,b}^\ell$ the activity of unit i on sample b , and $\mathcal{A}^\ell(X) = \{i : \exists b, a_{i,b}^\ell \neq 0\}$ the awake set. During the same forward step, a replay micro-batch drawn from the buffer is applied through the synapse mask

$$M_{ij}^\ell = \mathbf{1}[j \notin \mathcal{A}^{\ell-1}(X) \vee i \notin \mathcal{A}^\ell(X)], \quad (3)$$

i.e. iff the presynaptic unit is silent for the current input or the postsynaptic unit is asleep. The mask applies to the *entire* replay-induced parameter change, not only to the data gradient: with heavy-ball velocity v ($\mu = 0.9$), the local update direction $G = \varepsilon^\ell (a^{\ell-1})^\top$ and the per-step decay coefficient λ , a replay step is

$$v \leftarrow M \odot (\mu v + G) + (1 - M) \odot v, \quad W \leftarrow W + \eta M \odot v - \lambda M \odot W, \quad (4)$$

so the velocity is frozen (not zeroed) outside the mask, the decay is confined to it, and biases follow the unit mask $\mathbf{1}[i \notin \mathcal{A}^\ell(X)]$; unmasked momentum would re-inject the update into awake synapses [McKee et al., 2026, Bricken et al., 2023]. The readout row and its bias stay plastic, so old classes remain calibrated; the guarantees below concern the hidden computation. Within one step the order is: forward pass under the current suppression, unmasked waking update, awake sets from that pass, masked replay on the updated weights; replay is inferred with the suppression removed (a refractory unit may fire on replay and learn from it) and a burst reuses one mask (Algorithm 1).

Proposition 1 (Pre-silent updates are exactly invisible). *Fix X and let $\widehat{W}^\ell = W^\ell + \Delta^\ell$ with $\Delta_{ij}^\ell = 0$ whenever $j \in \mathcal{A}^{\ell-1}(X)$. Then every hidden activity on every sample of X is unchanged for any magnitude of Δ ; if the readout row satisfies the same condition, or is held fixed, the network output is unchanged as well.*

Proposition 2 (Post-asleep updates are invisible under a margin). *Let Δ^ℓ and Δb^ℓ additionally move synapses and biases of units $i \notin \mathcal{A}^\ell(X)$ with active presynaptic partners, and write $\delta_{i,b} = \sum_j \Delta_{ij}^\ell a_{j,b}^{\ell-1} + \Delta b_i^\ell$. If for every such unit and every sample $\sigma(z_{i,b}^\ell + \delta_{i,b}) < \tau_{k,b}^\ell$, where $\tau_{k,b}^\ell$ is the k -th largest value entering Φ_k on sample b (strict, so a tie cannot promote the unit; when fewer than k entries are positive, $\tau_{k,b}^\ell = 0$ and the condition reads $\sigma(\cdot) = 0$), the hidden activities on X are unchanged.*

Corollary 1 (Suppressed units are exact for any magnitude). *If $s_i^\ell = 1$ in (1), $a_{i,b}^\ell = 0$ for every sample and Proposition 2 holds with no margin condition.*

Proofs are in Appendix A. The default system does not enforce the margin, so we measured it over all 18,760 isolated replay updates of the record run (three seeds): the update altered the top hidden code of 0.32% of waking samples (0.29–0.36) and a waking prediction for 0.33% (almost all through the plastic readout; 0.001% with it isolated), violating the margin for $\sim 10^{-4}$ of asleep units per step, with a mean maximal logit change of 0.02 (0.001 isolated) on the one-hot scale. The isolation is exact for silent-presynaptic and suppressed units and near-exact elsewhere for the hidden code of the *current* batch; later inputs are the ablation’s business (Table 1), and Corollary 1 is why rotation makes the isolation robust.

Refractory rotation. After batch X_t , set $s_i^\ell(t+1) = \mathbf{1}[i \in \mathcal{A}^\ell(X_t)]$: every unit that fired sits out the next competition and is therefore asleep — consolidable, and exactly so — for X_{t+1} . Rotation widens the isolated channel $\rho_\ell = |\tilde{\mathcal{A}}^\ell \setminus \mathcal{A}^\ell(X)|/|\tilde{\mathcal{A}}^\ell|$ — the fraction of the units $\tilde{\mathcal{A}}^\ell$ activated by the replay micro-batch that are asleep for the current input — from 0.2–0.3 to ≈ 0.53 , at the price of running inference on the second-best coalition.

Internal control (Figure 1b). *When to replay:* each unit integrates its own use, $S_i \leftarrow S_i + \bar{a}_i(X_t)$ with $\bar{a}_i(X_t)$ the fraction of the batch’s samples on which it fired — a unit-level Process S [Borbély, 1982] — and a burst of n_b replay micro-batches fires when the mean pressure over the currently asleep units of all hidden layers exceeds θ ; consolidation resets the pressure of the units that received it ($S_i \leftarrow 0$). *When to rotate:* with fast and slow exponential averages of the batch-mean squared output error e_t , $\bar{e}_f$ ($\alpha_f=0.1$) and $\bar{e}_s$ ($\alpha_s=0.002$), rotation runs while $\bar{e}_f \geq \beta \bar{e}_s$ ($\beta \approx 1.1$ – 1.5) or while $\bar{e}_f \geq \gamma e_0$ with e_0 the chance-level error of the first 20 batches ($\gamma \approx 0.5$): the stream is new relative to the learner’s own recent history, or is not yet mastered. The detector is a standard drift monitor [Ross et al., 2012], loosely inspired by cholinergic novelty signalling [Hasselmo, 2006]; what is new is what it gates, and a single absolute threshold cannot serve across the regimes considered here (a converged ten-class i.i.d. error exceeds any level a two-class task dips under). At the activity fraction used throughout the gate is nearly neutral, and the default system runs rotation always on (Section 4).

Protocol. Split-MNIST (five tasks of two classes, five epochs per task) and split CIFAR-10 (same protocol, 32×32 luminance), each with its i.i.d. counterpart; a third regime feeds CIFAR-10 through a frozen V1-like patch front-end (1024-D, linear probe $\approx 49\%$); split CIFAR-100 appears in Section 5 for mechanism only. Architecture d –512–256– C , k -WTA 10%, 30% distance-dependent connectivity, 80% excitatory units; waking batches of 16; buffer $K=1000$ ($K=200$ in Appendix G). Baselines: BP with experience replay (BP+ER) at equal K ; the offline “night” (20 replay batches of 256 at $3 \times$ learning rate after each waking epoch, the best offline schedule in our sweep) on its preferred narrower core and on our substrate; interleaved ER for the local learner; no buffer. Three seeds (six




Figure 1: **Mechanism.** (a) During each waking batch, replay from the episodic buffer is written only into hidden synapses invisible to the current input — exactly so for silent presynaptic and suppressed units; refractory rotation forces recently used units into the consolidable set. (b) Unit-level homeostatic pressure triggers replay bursts; a relative-novelty gate decides when rotation runs. The default system runs rotation always on with one micro-batch per waking batch; the pressure trigger and the novelty gate are schedulers for sparse budgets and lower activity fractions.

on headline rows), mean $\pm$ s.d.; forgetting F is the mean drop of each earlier task’s accuracy from its own end to the end of training (per-task matrix in Appendix K). Evaluation runs with the suppression mask off (k -WTA only), so test order and batch size cannot affect it. Energy is an idealised 45-nm FP32 arithmetic proxy [Horowitz, 2014] (0.9 pJ per accumulate, one per synaptic event, against 4.6 pJ per multiply-accumulate) after rate-to-spike conversion with training-calibrated thresholds [Rueckauer et al., 2017], excluding data movement and membrane updates. Every development choice used the official test sets; Appendix K replicates the headline orderings on a held-out tenth of the training data ($\rho = 0.97$, mean gap 0.9); further protocol details are in Appendix D.

4 Results

Local sleep replaces the night. On split-MNIST at $K=1000$, refractory local sleep with one isolated replay micro-batch beside every waking batch reaches $92.7 \pm 0.3\%$ (six seeds) with no offline phase — above the best offline night ($90.7 \pm 0.4\%$ on its preferred narrow substrate; 90.9 ± 3.2 , highly variable across seeds, on ours) and far above unmasked interleaved ER at the same budget (89.7 ± 1.1) and BP+ER (89.3 ± 0.8). A night added on top changes nothing (92.7 ± 0.4). The headline configuration replays $2.4\times$ the night’s samples; within $1.2\times$ of the night’s samples (22% more, in $39\times$ more updates of 16) the system still reaches 91.7–91.8 (eight-sample micro-batches every batch, or 16-sample ones every second batch; Figure 2, Appendix D).

The same ordering on a split never used for selection. Every configuration in this paper was chosen on the official test set and frozen before this check, in which a stratified tenth of the training set is held out as the evaluation set and never trained on (during development those samples were training data only; no selection ever looked at them). Over six seeds, rotation reaches 91.6 ± 0.3 against the same-substrate night’s 89.0 ± 3.5 (paired margin $+2.6$ [$+0.2, +5.2$]), BP+ER’s 88.8 ± 0.3 ($+2.8$) and unmasked ER’s 85.4 ± 3.5 ($+6.3$); held-out numbers sit 0.9 points below test on average and the sequential ranking correlates at Spearman 0.97 (Appendix K, where the CIFAR-10 orderings replicate as well). Every other number in this paper is a development-phase result on the official test set.

What rotation and isolation each buy. Figure 2 (left): the full system is flat in replay batch size from 256 down to 8 ($92.3 \pm 0.1 \rightarrow 91.8 \pm 0.3$); at batch 16 unmasked replay without rotation sits at 87.0 ± 1.1 , isolation without rotation at 88.8 ± 0.8 , and the offline night falls from 90.9 to 87.3 ± 3.2 when its batches shrink to 16. The fourth cell — rotation with unmasked replay — reaches 92.1 ± 0.5 at batch 16, 91.2 ± 0.4 at 8 and 89.5 ± 0.9 at 4, where the masked system turns bimodal (63.4 ± 46.4): the factors are strongly sub-additive, rotation is the larger one ($+5.1$ alone, $+3.9$ on top of isolation), and isolation’s marginal accuracy is 0.6 at batches 8–16, within seed noise, with no static difference (94.7 either way). What isolation buys is the guarantee that no replay update touches the units the present input uses; without it the outcome rests on the rotation’s tolerance of the perturbation.

Component prices (Table 1). Removing rotation costs 3.9 sequential points (88.8 ± 0.8 : the same mask on natural silence alone), removing both 5.7 with a 3.7-fold larger seed deviation (87.0 ± 1.1), removing isolation alone 0.6 (above), and removing replay everything (19.6, the forgetting floor).




Figure 2: **Left (isolation):** split-MNIST accuracy against replay batch size, the number of updates fixed within each series (one per waking batch; 20 per epoch for the night, whose samples therefore fall with its batch): the full system is insensitive down to 8 samples; rotation with unmasked replay (dashed) trails it by 0.6–0.7 at 8–16 and holds at 4, where the masked system turns bimodal; without rotation, unmasked and isolated replay and the offline night all degrade. **Right (timing):** accuracy against the number of waking replay events: trickles fail, clocked bursts recover most of the gap, pressure-triggered bursts (star) dominate at sparse budgets, surprise triggering fires no events in this regime.

222 Replay confined to the readout row, hidden layers frozen for replay, is at chance (27.3 ± 15.1) and
 223 costs 12 static points: the effect is hidden-layer consolidation entirely. A random suppression of
 224 matched size recovers most of rotation’s gain (90.8 ± 1.3) but trails it by 1.9 points on both axes with
 225 $1.8\times$ the forgetting: alternating the coalition carries most of the effect, resting the recent winners the
 226 rest. The static column identifies who pays: rotation is the only statically costly component, about
 227 1.2 points at this activity fraction (94.7 against 95.9–96.1 without it, isolation being free either way).
 228 The unmasked control is not handicapped by the replay step (3η against a waking step of $\eta/16$): over
 229 replay gains $\{1/16, 1/4, 1, 3, 10\}$ both curves peak at the default (92.7 masked, 87.0 unmasked) and
 230 the unmasked control never closes the gap (Appendix I). *Mirror isolation* (waking updates barred
 231 from synapses that served the last replay batch, replay unmasked) reaches only 77.9 ± 1.6 : this
 232 reversed-mask control loses fifteen points, one instantiation of the protect-the-past direction rather
 233 than a verdict on that literature. *Soft rotation* (halving rather than silencing recent winners) loses
 234 on both axes ($90.2 \pm 0.4/86.4 \pm 0.3$): all-or-none OFF periods beat turned-down gain. *Optimiser*
 235 *state*: masked Adam trails heavy-ball SGD ($91.8 \pm 0.6/94.6 \pm 0.2$) because the guarantee freezes its
 236 moments outside each batch’s mask, where they go stale across a unit’s on/off cycles [Bricken et al.,
 237 2023]; letting the moments leak recovers most of the gap (92.6 ± 0.3) but surrenders the confinement
 238 of optimiser state (the leaked moments re-enter awake synapses at the next unmasked waking step),
 239 while a single velocity keeps both. Momentum-free SGD is no answer either: without the velocity’s
 240 noise averaging, isolated micro-batches collapse to chance unless the replay step is re-scaled by hand,
 241 and its best setting trails by 2.8/4.6; a two-sided learning-rate sweep shows these comparisons are
 242 plateaus, not knife-edges (Appendices I, H).

243 **When replay must be rare: homeostatic bursts, not surprise.** Figure 2 (right): at 12.5% of the
 244 replay events a clocked burst of 16 batches every 128 retains 87.3 ± 0.2 while an even trickle at the
 245 same budget collapses (81.9 ± 2.9). Unit-level pressure triggering does better still (89.1 ± 1.1 at a
 246 third of the events); surprise triggering fires no event, the converged error never crossing its threshold
 247 (19.0) — the opposite of the offline night, whose best trigger in our sweeps is novelty-timed.

248 **The i.i.d. price and the activity fraction.** Rotation’s static price is governed by the substrate’s
 249 activity fraction: 1.3–3.0 points at 5%, where gates were needed to refund it, about a point at 10%,
 250 where the gates are neutral (Table 1); on feature CIFAR-10 relaxing the fraction from 5 to 25%
 251 improves both axes monotonically ($40.5 \rightarrow 42.7 / 39.4 \rightarrow 45.6$), the local learner 1.0–1.3 points
 252 ahead of matched-sparsity BP at every level (Appendices E–F).

253 **Transfer to CIFAR-10.** On raw split CIFAR-10 the ordering transfers and widens: refractory local
 254 sleep $29.5 \pm 0.8 >$ night $26.2 \pm 0.6 >$ BP+ER $25.1 \pm 1.5 >$ unmasked ER $17.0 \pm 6.1 \approx$ no buffer

Table 1: **Component ablation** (512–256, k -WTA 10%, $K=1000$, one replay micro-batch of 16 per waking batch). Sequential = split-MNIST; static = i.i.d. MNIST. Headline rows six seeds, the rest three.

Variant	Sequential	i.i.d. (static)
full system: isolated replay + always-on rotation (heavy-ball SGD)	92.7 ± 0.3	94.7 ± 0.3
rotation gated by relative novelty, committed to bouts	92.3 ± 0.2	94.9 ± 0.3
rotation gated per batch	89.7 ± 1.4	95.3 ± 0.2
– rotation (same mask, Eq. 3, natural silence only)	88.8 ± 0.8	95.9 ± 0.2
– isolation, – rotation (unmasked interleaved replay)	87.0 ± 1.1	96.1 ± 0.1
– isolation, + rotation (rotation alone)	92.1 ± 0.5	94.7 ± 0.1
readout-only replay (hidden layers take no replay update)	27.3 ± 15.1	83.1 ± 2.4
random rotation (matched suppression count)	90.8 ± 1.3	92.8 ± 0.5
– replay (no buffer)	19.6 ± 0.0	95.8 ± 0.2
masked Adam, moments confined to the mask	91.8 ± 0.6	94.6 ± 0.2
masked Adam, moments leaking outside the mask	92.6 ± 0.3	94.3 ± 0.3
momentum 0 (pure SGD; η and replay gain re-tuned, else chance)	89.9 ± 0.3	90.1 ± 0.2
offline night instead of concurrent replay, same substrate	90.9 ± 3.2	—
the same night on its preferred narrow substrate	90.7 ± 0.4	—
mirror direction: protect the past	77.9 ± 1.6	92.3 ± 0.2
soft rotation: gain 0.5 instead of 0	90.2 ± 0.4	86.4 ± 0.3
unmasked ER at the same budget (Adam)	89.7 ± 1.1	95.4 ± 0.3
backprop (+ER sequential; plain static), dense	89.3 ± 0.8	97.5 ± 0.0

16.4. On the feature front-end the local ordering persists (refractory $41.5 \pm 1.1 > \text{night } 35.0 \pm 0.5 \gg$ unmasked 20.1 ± 7.3) while BP+ER reaches 43.2 ± 1.0 at its best learning rate — a 1.7-point lead. Transplanting the local network’s ingredients into the BP net one at a time (each at its better of two learning rates) attributes that gap to activation sparsity (width 44.5 ± 1.0 , +30% wiring 43.3 ± 0.9 , + k -WTA 40.5 ± 0.7 : at or below the local learner on the identical substrate). Converted to spiking inference ($T_s=8$) the default system runs at 93.8% for a proxy of 119 nJ per sample, $20.7\times$ below dense rate inference (Appendix D).

5 Scaling the mechanism: self-referenced decay and depth

The last tuned constant. Every result above runs the rule with a decay λ tuned once per dataset (10^{-3} on the image benchmarks). Split CIFAR-100 exposes what that constant hides: each class receives a tenth of the supervision events while the decay acts on every synapse at every step, so at the MNIST-tuned value the hidden weights follow the pure-decay trajectory and the network lands at chance (2.4% i.i.d. against 20–23% on a ten-class subset): the class count is a hidden variable every fixed decay is silently tuned against. We replace the constant by a per-layer ratio controller,

$$\lambda_\ell = \min\left(\lambda_{\max}, \rho \text{EMA}[u^\ell] / (\overline{|W^\ell|} + \epsilon)\right), \quad \rho = 0.25, \lambda_{\max} = 0.02, \quad (5)$$

where u^ℓ is the mean magnitude of the post-optimiser, *pre-decay* waking increment of layer ℓ (the decay never feeds its own drive; EMA coefficient 0.02 over unmasked waking steps, $\epsilon = 10^{-12}$), and W^ℓ , $B^{\ell-1}$ share λ_ℓ : the decay’s mean magnitude $\lambda_\ell \overline{|W^\ell|}$ is held at the fraction ρ of the mean data-driven update. ρ is the only control target; $\lambda_{\max}$, the EMA coefficient and ϵ are global safeguards never tuned per dataset. Isolation and rotation decide *where* and *when* to consolidate, the controller *how strongly* — by the principle behind every control signal here: novelty as a fast-over-slow error ratio, homeostasis as a unit’s own use, decay as a drive-over-weight ratio, all self-referenced, none absolute.

One ratio target, every dataset. With the same ρ on every dataset (Table 4 in Appendix J; e.g. split-MNIST $92.7/94.7 \rightarrow 91.7/96.1$, CIFAR-10 $29.5/27.6 \rightarrow 26.5/29.9$, sequential/static) the controller concedes 1–3 sequential points to the per-dataset tuned value and gains 0.5–3.6 on every static axis on the official test sets — the fixed decay had under-regularised every thick-signal regime — with 96.1 ± 0.2 the best i.i.d. accuracy we have seen from the two-layer substrate, a margin that did not replicate cleanly on the held-out split (Appendix K) and is therefore provisional. On CIFAR-100 it




Figure 3: **Depth is the same starvation on another axis, and two mechanisms cross it.** Sequential (left) and static (middle) accuracy against hidden depth (three seeds): the tested fixed decay collapses, the controller degrades gracefully, and controller + locally learned skip synapses hold the five-layer stack at the two-layer level. Right: the controller’s realised per-layer decay.

lifts the system 2.6–5 $\times$ over the manual rescue with no dataset-specific decay coefficient. The realised λ_ℓ self-differentiates as the tuning history would predict (10⁻⁵ on MNIST streams to 3 $\times$ 10⁻³ on raw CIFAR): the per-dataset ladder was a signal-to-decay ratio in disguise (Appendix J).

Depth (Figure 3). Under the tuned fixed decay without skips the static axis dies at four hidden layers (9.7 $\pm$ 0.7%) and the sequential axis turns bimodal at five (55.5 $\pm$ 39.3); under the controller every cell stays alive (90.6 $\rightarrow$ 90.0 $\rightarrow$ 87.6 sequential, 95.6 $\rightarrow$ 95.3 $\rightarrow$ 91.8 static). The residual toll is architectural: ResNet-style skip synapses — S^ℓ carrying $a^{\ell-2}$ into each deep hidden layer, learned by the same local product with a Kolen–Pollack feedback twin and the shared decay, replay-masked under the same pre-or-post-asleep rule (Appendix B) — restore the thick error path, and the five-hidden-layer system returns to the two-layer level: 91.5 $\pm$ 0.4/96.2 $\pm$ 0.3 against the two-layer controller cell’s 91.7/96.1, still 1.2 sequential points below the two-layer fixed-decay default. The bypass alone also rescues the fixed decay (91.3/93.6 at depth five), consistent with error-path attenuation as a principal cause of the collapse; the controller adds its static margin on top (the fixed decay was not re-tuned per depth; the claim is that the controller needs no re-tuning). BP+ER at the same widths is depth-insensitive ($\approx$ 88–89) but forgets 2–3 $\times$ more; the local system with skips and controller beats it at every depth tested. On CIFAR-100, by contrast, depth is a net loss that skips only half refund, while width pays (Appendix J).

6 Discussion

What this offers agents. The construction is a recipe: if an agent’s representation is sparse enough that a batch leaves most units silent, and its learning rule yields a local error signal, consolidation can run in the silent degrees of freedom between the steps of the stream, with no separate offline phase — in the supervised setting studied here, with a guarantee on the computation in use and a replay budget an online system can afford. The invariance does not protect a finished forward pass: it closes the channel through which consolidation would perturb the units the present stream is using. Its accuracy value is large without rotation (−5.7 points) and small with it (−0.6): once the coalition alternates, the live units tolerate unmasked replay, and what isolation adds is the guarantee that they need not. We treat the biology as motivating analogy, not mechanistic equivalence; the refractory rule mirrors the ON/OFF finding of Driessen et al. [2026], and the trigger dissociation (homeostatic during wake, novelty-timed offline) is a testable prediction.

Limitations. All results are at MLP scale on MNIST/CIFAR variants with three to six seeds; every configuration choice during development was made on the official test sets, and the held-out check covers one split of MNIST and CIFAR-10; the protocol is five epochs per task, not a single pass (a single pass keeps the ordering, Appendix I). BP+ER (an external reference, not a state-of-the-art claim; DER++-style baselines are not included) keeps a 1.7-point lead on feature-based CIFAR-10 (0.5 held-out), attributed by the tested decomposition to dense activation; rotation keeps a static price of about a point; CIFAR-100 is floor-regime for every method; the energy figure is an arithmetic proxy. Within these limits the conclusion stands: consolidation does not require a separate offline phase — the learner consolidates between the steps of its stream, above the strongest offline-rehearsal schedule at 2.4 $\times$ its replay samples and within a point of it at 1.2 $\times$; code will be released.

References

- Abbasi, A., Nooralinejad, P., Braverman, V., Pirsiavash, H., Kolouri, S. (2022). Sparsity and heterogeneous dropout for continual learning in the null space of neural activations. In *Conference on Lifelong Learning Agents (CoLLAs)*, PMLR 199.
- Akrout, M., Wilson, C., Humphreys, P., Lillicrap, T., Tweed, D. (2019). Deep learning without weight transport. In *NeurIPS* 32.
- Arani, E., Sarfraz, F., Zonooz, B. (2022). Learning fast, learning slow: A general continual learning method based on complementary learning system. In *ICLR*.
- Bazhenov, A., Delanois, J.E., Krishnan, G.P. (2026). Not just after one: Sleep-inspired replay prevents catastrophic forgetting after sequential tasks. *arXiv:2606.08447*.
- Benna, M.K., Fusi, S. (2016). Computational principles of synaptic memory consolidation. *Nature Neuroscience*, 19:1697–1706.
- Borbély, A.A. (1982). A two process model of sleep regulation. *Human Neurobiology*, 1:195–204.
- Bricken, T., Davies, X., Singh, D., Krotov, D., Kreiman, G. (2023). Sparse distributed memory is a continual learner. In *ICLR*.
- Driessen, K., Squarcio, F., Tononi, G., Cirelli, C. (2026). Induction of cortical ON/OFF periods in awake mice fulfills sleep functions. *Nature Neuroscience*, 29:1954–1965.
- Farajtabar, M., Azizan, N., Mott, A., Li, A. (2020). Orthogonal gradient descent for continual learning. In *AISTATS*, PMLR 108.
- Golden, R., Delanois, J.E., Sanda, P., Bazhenov, M. (2022). Sleep prevents catastrophic forgetting in spiking neural networks by forming a joint synaptic weight representation. *PLoS Computational Biology*, 18:e1010628.
- Golkar, S., Kagan, M., Cho, K. (2019). Continual learning via neural pruning. *arXiv:1903.04476*.
- Harun, M.Y., Gallardo, J., Hayes, T.L., Kemker, R., Kanan, C. (2023). SIESTA: Efficient online continual learning with sleep. *TMLR*.
- Hasselmo, M.E. (2006). The role of acetylcholine in learning and memory. *Current Opinion in Neurobiology*, 16:710–715.
- Horowitz, M. (2014). Computing’s energy problem (and what we can do about it). In *IEEE International Solid-State Circuits Conference (ISSCC)*, pp. 10–14.
- Iyer, A., Grewal, K., Velu, A., Souza, L.O., Forest, J., Ahmad, S. (2022). Avoiding catastrophe: Active dendrites enable multi-task learning in dynamic environments. *Frontiers in Neuroinformatics*, 16:846219.
- Kaski, S., Kohonen, T. (1994). Winner-take-all networks for physiological models of competitive learning. *Neural Networks*, 7:973–984.
- Klasson, M., Kjellström, H., Zhang, C. (2023). Learn the time to learn: Replay scheduling in continual learning. *TMLR*.
- Kolen, J.F., Pollack, J.B. (1994). Backpropagation without weight transport. In *IEEE International Conference on Neural Networks (ICNN)*, pp. 1375–1380.
- Kubo, Y., Delanois, J.E., Bazhenov, M. (2025). Toward lifelong learning in equilibrium propagation: Sleep-like and awake rehearsal for enhanced stability. *arXiv:2508.14081*.
- Lin, E., Luo, W., Jia, W., Yang, S., Xu, W. (2026). Online continual learning via spiking neural networks with sleep enhanced latent replay. In *14th International Conference on Intelligent Control and Information Processing (ICICIP)*, pp. 360–367.
- Lässig, F., Aceituno, P.V., Sorbaro, M., Grewe, B.F. (2023). Bio-inspired, task-free continual learning through activity regularization. *Biological Cybernetics*, 117:345–361.
- Maeda, M., Miyajima, H. (1999). Competitive learning methods with refractory and creative approaches. *IEICE Transactions on Fundamentals*, E82-A(9):1825–1833.
- Mallya, A., Lazebnik, S. (2018). PackNet: Adding multiple tasks to a single network by iterative pruning. In *CVPR*.

369 Masse, N.Y., Grant, G.D., Freedman, D.J. (2018). Alleviating catastrophic forgetting using context-dependent
370 gating and synaptic stabilization. *PNAS*, 115:E10467–E10475.

371 McKee, K., Hazy, T., Zheng, Y., Bugaud, Z., Miconi, T. (2026). Cortex-inspired continual learning: Unsupervised
372 instantiation and recovery of functional task networks. *arXiv:2604.24637*.

373 Payeur, A., Guerguiev, J., Zenke, F., Richards, B.A., Naud, R. (2021). Burst-dependent synaptic plasticity can
374 coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24:1010–1019.

375 Pham, Q., Liu, C., Hoi, S. (2021). DualNet: Continual learning, fast and slow. In *NeurIPS* 34.

376 Robinson, B.S., Lau, C.W., New, A., Nichols, S.M., Johnson, E.C., Wolmetz, M., Coon, W.G. (2022).
377 Continual learning benefits from multiple sleep mechanisms: NREM, REM, and synaptic downscaling.
378 *arXiv:2209.05245*.

379 Ross, G.J., Adams, N.M., Tasoulis, D.K., Hand, D.J. (2012). Exponentially weighted moving average charts for
380 detecting concept drift. *Pattern Recognition Letters*, 33:191–198.

381 Rueckauer, B., Lungu, I.-A., Hu, Y., Pfeiffer, M., Liu, S.-C. (2017). Conversion of continuous-valued deep
382 networks to efficient event-driven networks for image classification. *Frontiers in Neuroscience*, 11:682.

383 Saha, G., Garg, I., Roy, K. (2021). Gradient projection memory for continual learning. In *ICLR*.

384 Saper, C.B., Scammell, T.E., Lu, J. (2005). Hypothalamic regulation of sleep and circadian rhythms. *Nature*,
385 437:1257–1263.

386 Shen, J., Ni, W., Xu, Q., Tang, H. (2024). Efficient spiking neural networks with sparse selective activation for
387 continual learning. In *AAAI* 38, pp. 611–619.

388 Sorrenti, A., Bellitto, G., Proietto Salanitri, F., Pennisi, M., Palazzo, S., Spampinato, C. (2025). Wake–sleep
389 consolidated learning. *IEEE Transactions on Neural Networks and Learning Systems*, 36(7):12668–12679.

390 Tadros, T., Krishnan, G.P., Ramyaa, R., Bazhenov, M. (2022). Sleep-like unsupervised replay reduces catastrophic
391 forgetting in artificial neural networks. *Nature Communications*, 13:7742.

392 Tilley, M.J., Miller, M., Freedman, D.J. (2023). Artificial neuronal ensembles with learned context dependent
393 gating. In *ICLR*.

394 Tononi, G., Cirelli, C. (2014). Sleep and the price of plasticity: From synaptic and cellular homeostasis to
395 memory consolidation and integration. *Neuron*, 81:12–34.

396 Vitter, J.S. (1985). Random sampling with a reservoir. *ACM Transactions on Mathematical Software*, 11:37–57.

397 Vyazovskiy, V.V., Olcese, U., Hanlon, E.C., Nir, Y., Cirelli, C., Tononi, G. (2011). Local sleep in awake rats.
398 *Nature*, 472:443–447.

399 Wang, S., Li, X., Sun, J., Xu, Z. (2021). Training networks in null space of feature covariance for continual
400 learning. In *CVPR*.

401 Whittington, J.C.R., Bogacz, R. (2017). An approximation of the error backpropagation algorithm in a predictive
402 coding network with local Hebbian synaptic plasticity. *Neural Computation*, 29:1229–1262.

403 A Proofs

404 *Proof of Proposition 1.* By induction on ℓ , for every sample b . Assume $a_{i,b}^{\ell-1}$ is unchanged (true
405 for $\ell = 1$ since $a_{i,b}^0 = x_b$). For any unit i , $\hat{z}_{i,b}^\ell - z_{i,b}^\ell = \sum_j \Delta_{ij}^\ell a_{j,b}^{\ell-1}$, and each term vanishes: if
406 $j \in \mathcal{A}^{\ell-1}(X)$ then $\Delta_{ij}^\ell = 0$, otherwise $a_{j,b}^{\ell-1} = 0$ for every b . Hence $z_{i,b}^\ell$, and $a_{i,b}^\ell$ through (1), are
407 unchanged; the induction closes at the last hidden layer, and at the output layer whenever the readout
408 row satisfies the same condition. $\square$

409 *Proof of Proposition 2.* A unit outside $\mathcal{A}^\ell(X)$ contributes $a_{i,b}^\ell = 0$ on every sample. After the update
410 its pre-activation on sample b is $z_{i,b}^\ell + \delta_{i,b}$; under the stated margin it is still excluded by Φ_k (or
411 rectified to zero), so $a_{i,b}^\ell = 0$ still. The sample’s original winners keep their pre-activations — their
412 synapses from active presynaptic units lie outside Δ^ℓ by the mask, and their silent-presynaptic terms
413 vanish as in Proposition 1 — so the same k units win with the same values, every other unit’s input is
414 untouched, and the layer’s activities are identical. When fewer than k units are positive on sample b ,
415 $\tau_{k,b}^\ell = 0$ and the condition reads $\sigma(z_{i,b}^\ell + \delta_{i,b}) = 0$: a zero-valued entry carries no activity whether or
416 not Φ_k selects it. The readout row is excluded from both propositions in the default system because it
417 is deliberately kept plastic. $\square$

418 B One training step, the three channels and the skip path

419 Algorithm 1 lists one waking step of the default system with the state each quantity is read from.
420 Two conventions matter for the guarantees. First, the awake sets are read from the waking forward
421 pass, i.e. from the weights before the unmasked waking update, while the masked replay step acts
422 on the updated weights; Propositions 1–2 hold verbatim for a mask computed on the weights being
423 updated, and the drift reported in Section 3 is measured under the implemented order. Second,
424 the replay micro-batch is inferred with the suppression removed, so a refractory unit may fire on
425 replay and receive the update (Corollary 1 makes that update invisible to the waking batch for any
426 magnitude), and a burst of n_b micro-batches reuses the mask of its waking batch. The waking step is
427 $\eta_w = \eta m/256$ for a waking batch of m samples (the per-epoch drift of a batch-256 protocol), the
428 replay step $\eta_r = 3\eta$ per micro-batch, with $\eta = 0.02$ throughout.

Algorithm 1 One waking step of the default system (rotation always on, one burst per batch).

Require: weights W_t , velocities v_t , suppression $s(t)$, buffer $\mathcal{M}$, pressures S

- 1: $a \leftarrow \text{forward}(X_t; W_t, s(t))$, $\varepsilon \leftarrow \text{sweep}(a, y_t)$ Eqs. (1)–(2)
- 2: $(W, v) \leftarrow \text{unmasked heavy-ball step on } (W_t, v_t) \text{ with } G^\ell = \varepsilon^\ell (a^{\ell-1})^\top$, rate η_w wake update
- 3: $\mathcal{A}^\ell \leftarrow \{i : \exists b, a_{i,b}^\ell \neq 0\}$ for every hidden layer; $S_i \leftarrow S_i + \bar{a}_i(X_t)$ from the pre-update a
- 4: $M^\ell \leftarrow \text{Eq. (3)}$, unit mask $\mathbf{1}[i \notin \mathcal{A}^\ell]$, readout row and bias unmasked
- 5: $s_i^\ell(t+1) \leftarrow \mathbf{1}[i \in \mathcal{A}^\ell]$ refractory rotation
- 6: write (X_t, y_t) into $\mathcal{M}$ (reservoir)
- 7: **for** each of the n_b replay micro-batches $(X_r, y_r) \sim \mathcal{M}$ **do**
- 8: $a_r \leftarrow \text{forward}(X_r; W, s=0)$, $\varepsilon_r \leftarrow \text{sweep}(a_r, y_r)$ suppression removed
- 9: $(W, v) \leftarrow \text{masked step, Eq. (4), with } G^\ell = \varepsilon_r^\ell (a_r^{\ell-1})^\top$, rate η_r , masks M^ℓ
- 10: **end for**
- 11: $S_i \leftarrow 0$ for every $i \notin \mathcal{A}^\ell$ if a burst ran discharge
- 12: **return** $W_{t+1} = W$, $v_{t+1} = v$, $s(t+1)$

429 **The three channels.** Table 2 lists the three ways a synapse enters the mask (3) and what each
430 guarantees; the rotation of Section 3 exists to move synapses from the second row into the third.

431 **A numerical example.** On one sample with $k = 2$ and $\sigma(z) = (3, 2, 0.5, 0)$, units 3 and 4 are
432 asleep and $\tau_k = 2$. Any replay update with $\sigma(z_3 + \delta_3) < 2$ and $\sigma(z_4 + \delta_4) < 2$ leaves the activity
433 vector $(3, 2, 0, 0)$ unchanged, because the two winners’ own inputs are not touched; if unit 4 is
434 refractory ($s_4 = 1$) its bound is void and δ_4 may be arbitrarily large; and every synapse leaving unit 4
435 into the next layer may move by any amount, since its presynaptic activity is zero.

Channel	Condition	Invisible to X	Source
presynaptic silent	$j \notin \mathcal{A}^{\ell-1}(X)$	for any magnitude	Prop. 1
postsynaptic naturally asleep, active pre	$i \notin \mathcal{A}^{\ell}(X), s_i^{\ell} = 0$	under the margin $\tau_{k,b}^{\ell}$	Prop. 2
postsynaptic suppressed	$s_i^{\ell} = 1$	for any magnitude	Cor. 1

Table 2: **The three consolidation channels** opened by the synapse mask.

Skip synapses. For $\ell \geq 3$ each deep hidden layer receives a bypass, $z^{\ell} = W^{\ell} a^{\ell-1} + S^{\ell} a^{\ell-2} + b^{\ell}$, with S^{ℓ} under the same Dale projection and a feedback twin $B_S^{\ell} \in \mathbb{R}^{n_{\ell+2} \times n_{\ell}}$ that carries the bypass target’s burst back to its source,

$$\varepsilon^{\ell} \leftarrow \varepsilon^{\ell} + \sigma'(z^{\ell}) \odot \left((B_S^{\ell})^{\top} [\kappa \tanh(\varepsilon^{\ell+2}/\kappa) \odot \mathbf{1}(a^{\ell+2} > 0)] \right),$$

S^{ℓ} and $B_S^{\ell-2}$ learned by the local product of the main text with the target layer’s decay λ_{ℓ} , and replay-masked by the same rule two layers apart, $M_{S,ij}^{\ell} = \mathbf{1}[j \notin \mathcal{A}^{\ell-2}(X) \vee i \notin \mathcal{A}^{\ell}(X)]$, so that Propositions 1–2 extend unchanged.

C The price of the substrate

Table 3 traces the path from a dense backpropagation MLP to the substrate used throughout, on i.i.d. MNIST. The cumulative ladder prices the entire constraint set at 1.2 points below the dense reference for $4\times$ fewer synaptic operations per pass; the local burst-error objective itself costs nothing. Settling cannot merely be shortened ($T=1$ collapses); the single-phase burst sweep matches five-step settling in one pass. The Kolen–Pollack decay and even restored weight transport are accuracy-neutral at this scale — the decay earns its keep as a regulariser on harder inputs. Spiking-flavoured training fails outright: spiking is affordable at inference, not as a training-time code. On the plain substrate SGD loses 1.3 points to Adam, yet inside the consolidation loop it wins on both axes: per-synapse adaptive state is an asset while every synapse trains on every batch and a liability once updates become mask-confined.

Table 3: **Substrate ablation** on i.i.d. MNIST (256–128 hidden core, three seeds). “Syn. ops” is the fraction of the dense MLP’s synaptic operations per forward pass. Upper block cumulative; lower block replaces one ingredient at its stage.

Substrate	Accuracy	Syn. ops
dense backpropagation MLP (autograd reference)	97.7 ± 0.1	1.00
local burst-error learner, unconstrained	97.8 ± 0.1	0.85
+ bounded error, Dale’s law, sign-concordant feedback (no transport)	97.1 ± 0.0	0.85
+ k -WTA 10%	97.0 ± 0.2	0.80
+ relaxation $T=5$	97.0 ± 0.1	0.80
+ 30% distance-dependent connectivity	96.5 ± 0.2	0.24
+ burst gate and burst baseline	96.3 ± 0.2	0.24
+ single-phase burst sweep replacing relaxation (<i>final substrate</i>)	96.5 ± 0.1	0.24
random wiring instead of distance-dependent, equal density	95.4 ± 0.0	0.24
one settling step ($T=1$) instead of five	17.9 ± 6.8	0.80
weight transport restored (mirroring, no KP decay)	96.3 ± 0.2	0.24
Kolen–Pollack decay removed	96.3 ± 0.1	0.24
heavy-ball SGD instead of Adam ($\eta = 0.02$)	95.2 ± 0.3	0.24
k -WTA 5% / 2% instead of 10%	$95.9 \pm 0.1 / 94.2 \pm 0.2$	0.79/0.78
multiplicative burst coding	78.8 ± 1.9	0.24
training on $T_s=8$ spike counts	80.0 ± 0.8	0.24

D Protocol details, cost accounting and the operating frontier

The V1-like front-end consists of 256 normalised random $6\times 6\times 3$ patches sampled from the training images, applied with stride 2, rectified against each filter’s mean response over training images, and

quadrant-average-pooled (1024-D); nothing in it is learned. Split CIFAR-10 uses 32×32 luminance so that distance-dependent wiring keeps its image geometry. The buffer uses reservoir writing; the night reference uses a 256–128 core, its preferred width in our sweeps. Runs are deterministic at a fixed CPU thread count.




Figure 4: **Training data flow.** Each waking batch is inferred under the current suppression mask; its awake sets define the synapse and unit masks for the replay micro-batch applied in the same step and the rotation mask for the next batch. The masked step confines the velocity and the decay to the mask; the readout row is the one deliberately plastic exception. Unit-level pressure decides when replay bursts run.

459

Energy. Converted to spiking inference ($T_s=8$), the default system runs at 93.8% (rate accuracy 94.7%) for a proxy of 119 nJ per sample against 2461 nJ for dense rate inference ($20.7\times$) and 526 nJ for event-driven rate inference ($4.4\times$); accuracy is monotone in T_s , whereas the rotation-free control degrades beyond $T_s=8$ ($94.6 \rightarrow 90.1$ at $T_s=32$), consistent with rotation shaping sparser codes.

Updates and wall-clock. Samples are not the only cost: the headline schedule makes 18,760 masked updates against the night’s 480 ($39\times$), and its training wall-clock on one CPU is 6.6 min against 0.6 for the night, 0.4 with no replay and 6.0 for unmasked interleaved replay — the per-batch replay step, not the sample count, is the online price of consolidating during the stream, in an unoptimised implementation whose mask construction is dense. Replay cost is reported in samples ($R = \text{replay batches} \times \text{batch size}$). The night costs $480 \times 256 \approx 1.2 \times 10^5$ samples per run. The headline configuration (one 16-sample micro-batch per waking batch) costs $18,760 \times 16 \approx 3.0 \times 10^5$ ($2.4\times$ the night) for 92.7 ± 0.3 . Two schedules meet the night’s budget within $1.2\times$: eight-sample micro-batches at every waking batch (1.5×10^5 , 91.8 ± 0.3) and 16-sample micro-batches at every second one (1.5×10^5 , 91.7 ± 1.0 sequential, 93.5 ± 0.4 static); pressure-triggered bursts with batch 64 cost 1.8×10^5 ($1.4\times$) for 89.1 ± 1.1 , and halving the cadence with doubled batches costs 1.3 points (91.4 ± 0.3). Accuracy is bought by *frequency* of small isolated updates rather than by replayed samples. Figure 5 draws the operating frontier of the record substrate.

Baseline tuning and paired comparison. The offline-night reference was swept over replay batch count (20, 40 per night), replay batch size (16, 64, 256), learning-rate gain ($3\times$), core width (256–128 and 512–256), buffer size (200, 1000, 5000) and timing (clocked after each epoch, novelty-timed, surprise-timed), and its best sequential cell is reported; BP+ER was swept over learning rate ($3 \cdot 10^{-4}$, 10^{-3} , $3 \cdot 10^{-3}$), width and buffer size; the local learner over η (Appendix H), replay batch size, cadence and the gates of Table 1. Selection everywhere was sequential accuracy at $K=1000$ with the static axis reported alongside, on the official test set (Appendix K). Pairing seeds by index, the local system’s margin over the narrow-substrate night is +2.1 points (bootstrap 95% CI [+1.4, +2.5], three seed pairs; Welch $p = 0.004$), over the same-substrate night +1.7 ([+0.1, +4.6], six pairs; the night’s seed variance makes this one inconclusive, $p = 0.25$), over unmasked interleaved ER +5.8 ([+5.0, +7.0]) and over BP+ER +3.5 ([+2.7, +4.2]).




Figure 5: **Operating frontier** at the 10% activity fraction under heavy-ball SGD: split-MNIST sequential accuracy against i.i.d. static accuracy. Blue: non-dominated configurations; purple: dominated; grey: the 5% family of earlier phases. No single configuration wins both axes; always-on rotation is the sequential default and the adaptive decay extends the static side.




Figure 6: **Left:** the substrate decomposition on feature CIFAR-10 — transplanting the local network’s ingredients into BP+ER one at a time (each cell at its better of two learning rates; six seeds except the local learner, three). **Right:** spiking inference, i.i.d. MNIST accuracy against the idealised 45-nm FP32 arithmetic-energy proxy for $T_s \in \{4, 8, 16, 32\}$, with the dense (square) and event-driven (diamond) rate references; rotation-trained networks are monotone in T_s , the rotation-free control degrades.

488 E Rotation gates at the 5% activity fraction

489 At the 5% fraction of the earlier phases, always-on rotation cost 1.3–3.0 points on i.i.d. MNIST
490 (95.0 $\rightarrow$ 92.0–93.7). Gating rotation by unit-level pressure only trades the axes (static 91.6 $\rightarrow$ 94.0
491 as sequential 88.8 $\rightarrow$ 80.7). The relative-novelty gate removes the price at no sequential cost on
492 MNIST (static 95.0 ± 0.4 with rotation on 1.4–17% of batches; sequential 90.3–90.9). On split
493 CIFAR-10 the price reverses — rotation gains +2.5 to +6.5 static points in the underfit regime —
494 while the pure novelty gate misfires there (sequential 19.1 ± 1.1 : fast/slow $\rightarrow$ 1 on a stream that
495 never becomes predictable), which the unmastered clause resolves with one setting across all four
496 regimes (27.2 ± 0.7 / 25.3 ± 1.0 on CIFAR, 90.4 ± 0.9 / 94.3 ± 0.2 on MNIST). Those gate numbers
497 use masked Adam; under heavy-ball SGD the per-batch gate’s flicker destroys the optimiser’s gain
498 and the gate must be committed to minimum-duration bouts ($M=2048$ batches, loosely modelled on
499 the consolidated bouts of the sleep–wake switch [Saper et al., 2005]; Table 1). Figure 7 shows the
500 three regimes.




Figure 7: **Rotation at 5%: price, trade-off and gate.** Left: on i.i.d. MNIST the novelty gate restores no-rotation accuracy. Middle: on split-MNIST rotation is required and the gate keeps it. Right: on split CIFAR-10 rotation is a gain on both axes.

F Probes of the residual static price

Bout length and the OFF-side refractory. Bout length is a proper knob: $M=512$ gives $91.7 \pm 0.4/94.0 \pm 0.5$, recovering most of what the per-batch gate loses, while $M=4096$ drives the static duty to 0.94 and the refund vanishes. Completing the switch with an OFF-side refractory (triggers ignored for M_{off} batches after a bout) is a clean null: short refractories never engage on the static stream, and a long one (4096) lowers the static duty to 0.53 without moving static accuracy (94.0 ± 0.7) while monotonically eroding the sequential axis ($92.1 \rightarrow 91.0$). The residual static price sits in the rotation of settled coalitions itself, not in how often the gate fires.

Structural exemption and synaptic anchoring. Sparing the top- q units by long-term use from ever being benched traces a smooth trade-off ($q=0.10$: $90.9 \pm 0.1/94.2 \pm 0.2$; $q=0.25$: $89.1 \pm 0.4/95.4 \pm 0.2$) that the bout gate dominates: temporal commitment beats structural exemption. Two-timescale synapses in the spirit of synaptic consolidation cascades [Benna and Fusi, 2016] — a slow anchor absorbing the fast weight at rate μ while the fast weight decays toward it at rate λ , ticking only on waking updates — do not dissolve the price (no cell exceeds the anchor-free static accuracy); weak symmetric coupling acts as a sequential stabiliser at 5% (92.4 ± 0.2 , variance halved) and is neutral at 10% ($92.7 \pm 0.4/94.6 \pm 0.3$).

G Small buffers and sleep-anchored down-selection

At $K=200$ the night appeared to keep a 2.6-point edge; on the same 512–256/5% substrate the night reaches 84.0 ± 3.4 while noisy pressure-burst local sleep reaches 84.3 ± 1.0 and local+night 84.8 ± 0.9 — parity, with three-fold smaller variance for the wake-side scheme. The strongest tiny-buffer number remains the narrower night (86.9), so the offline window retains an advantage only in combination with narrow substrates. Replay diversity (input noise $\sigma=0.2$) is a wake-side lever — it lifts every local scheme ($81.8 \pm 2.9 \rightarrow 84.3 \pm 1.0$) and hurts the night — and unmasked waking replay collapses (71.0 ± 1.7). Anchoring structural plasticity (prune the weakest live synapses, regrow with a distance bias) to consolidation windows, motivated by synaptic homeostasis [Tononi and Cirelli, 2014], is free on the continuous substrate (90.9 ± 0.2 vs. 90.8 ± 0.1), rescues the wide-sparse night (87.2 ± 3.8 vs. 84.1 ± 1.3), and harms the episodic-burst substrate (84.6 ± 0.5 vs. 89.9 ± 0.1): down-selection, like replay, needs either density or an offline window.

H Learning-rate sensitivity

At the 5% fraction, under always-on rotation, SGD holds a sequential plateau over $\eta \in [0.01, 0.02]$ ($91.9 \pm 0.5/92.0 \pm 0.6$) and a static plateau over $[0.02, 0.05]$ ($93.3-93.7$), decaying gracefully outside ($86.3/89.2$ at $\eta=0.1$); masked Adam, swept over a tenfold range, has its sequential optimum at the default (90.8 ± 0.1 at 10^{-3} ; 88.6 at $3 \cdot 10^{-4}$, 88.9 at $3 \cdot 10^{-3}$), and no Adam setting reaches the

heavy-ball joint point. Unmasked SGD is poor and unstable at every rate (84–88, seed deviations 1.8–4.1). The same two-sided sweep at the 10% record fraction (three seeds) gives the same verdict: heavy-ball SGD holds $92.2 \pm 0.7/93.3 \pm 0.2$ at $\eta=0.01$, $92.7 \pm 0.3/94.7 \pm 0.3$ at 0.02 (the joint optimum) and $91.5 \pm 0.6/94.7 \pm 0.3$ at 0.05, decaying at 0.005 ($91.3/91.3$) and destabilising at 0.1 (79.5 ± 17.3); masked Adam has its optimum at the default ($91.8 \pm 0.6/94.6 \pm 0.2$ at 10^{-3} ; $89.9/89.9$ at $3 \cdot 10^{-4}$, $90.9/94.3$ at $3 \cdot 10^{-3}$). The optimiser comparison of Table 1 is made at each optimiser’s best rate.

I Momentum, exactness and other negative results

Momentum. With momentum zero the default replay step (3η per micro-batch against the batch-scaled waking step $\eta m/256$, Appendix B) diverges or starves hidden activity at every learning rate tested ($\eta \in [0.02, 0.4]$: $\approx 10\%$, hidden dimensionality 0); the waking path alone learns normally at the matched effective step, so the collapse is the replay path’s: the velocity’s $\sim 10\times$ averaging of the 16-sample replay micro-batches is what heavy-ball provides. Re-scaling the replay gain ($\eta=0.2$, gain 0.3) makes pure SGD viable at $89.9 \pm 0.3/90.1 \pm 0.2$ ($91.7 \pm 0.2/91.4 \pm 0.0$ with the adaptive decay); other settings are fragile.

Exactness measurement. Under the 5% Adam-era configuration the isolated update altered the top hidden code of 1.0% of waking samples per step (against 0.32% for the record configuration); on raw CIFAR-10 (one seed) the hidden-code rate is 0.46% while the prediction rate rises to 3.6%, low-margin predictions flipping more easily through the plastic readout. The numbers in Section 3 are population rates over every replay step of a full run. They are taken under the implemented order of Algorithm 1, in which the awake sets come from the activities before the waking update; re-inferring the batch after that update and masking on the post-update awake sets lowers the hidden-code rate to 0.30% (0.17% with the readout isolated) and leaves the prediction rate (0.30%; 0.000% isolated), the margin-violation rate and the accuracy unchanged (92.2 vs. 92.0, one seed); the residual drift is the unenforced margin, not the ordering.

Replay gain. Sweeping the replay gain over $\{1/16, 1/4, 1, 3, 10\}$ (default 3; $\eta_r = 3\eta$ against a waking step of $\eta/16$) moves masked replay $71.9 \pm 6.9 \rightarrow 88.2 \pm 0.8 \rightarrow 91.8 \pm 0.4 \rightarrow 92.7 \pm 0.3 \rightarrow 90.8 \pm 0.9$ and unmasked replay without rotation $63.8 \pm 1.9 \rightarrow 83.3 \pm 4.3 \rightarrow 86.2 \pm 4.2 \rightarrow 87.0 \pm 1.1 \rightarrow 85.5 \pm 1.5$ (sequential, three seeds): both peak at the default, and no replay gain makes the unmasked control competitive.

Single-pass streams. With one epoch per task instead of five (everything else as in the headline cells, three seeds), refractory local sleep reaches 91.8 ± 0.4 ($F=5.2$) against 92.7 with five epochs, BP+ER 90.1 ± 0.7 , the offline night 75.4 ± 0.5 (one night per task is too few), unmasked interleaved replay 68.3 ± 19.1 and no buffer 18.8: the ordering of the five-epoch protocol survives a single pass, with the offline schedule the one that does not.

Falsified under matched budgets. A frozen dentate-gyrus-style sparse expansion front-end (input decorrelation does not propagate through learned k -WTA layers: input task overlap drops $0.79 \rightarrow 0.18$ yet every downstream metric worsens by 3–11 points); generative (REM) replay as a substitute for episodic samples; margin-based reverse learning; surprise-gated memory writing; multiplicative burst coding; absolute-threshold rotation gates; soft (gain-reduction) rotation; mirror-direction isolation; OFF-side gate refractories; strongly coupled synaptic anchors; utility-exempt rotation.

575 J The controller: composition rules, CIFAR-100 and reproducibility

Table 4: **One dataset-independent ratio target replaces per-dataset decay tuning.** Sequential / static accuracy under the tuned fixed decay (10^{-3} ; 10^{-4} with $3\times$ targets as the CIFAR-100 rescue) and the controller (5) at $\rho = 0.25$ everywhere (drive form; a gradient-referenced form behaves equivalently, Appendix J). Full system in every cell; CIFAR-100 numbers are floor-regime and reported for ordering only (BP+ER 6–10%).

	fixed λ (tuned)		controller	
	seq.	static	seq.	static
split-MNIST	92.7 $\pm$ 0.3	94.7 $\pm$ 0.3	91.7 $\pm$ 0.2	96.1 $\pm$ 0.2
split CIFAR-10	29.5 $\pm$ 0.8	27.6 $\pm$ 0.4	26.5 $\pm$ 1.3	29.9 $\pm$ 0.3
feature CIFAR-10	41.5 $\pm$ 1.1	41.8 $\pm$ 1.1	38.9 $\pm$ 2.0	43.0 $\pm$ 1.2
split CIFAR-100	1.4 $\pm$ 0.5	1.1 $\pm$ 0.1	3.4 $\pm$ 0.6	3.7 $\pm$ 0.3
+ features	1.2 $\pm$ 0.3	1.2 $\pm$ 0.2	5.7 $\pm$ 0.8	6.0 $\pm$ 0.7

576 A gradient-referenced form of the controller ($\rho \eta \text{EMA}[\overline{[g^\ell]}] / \overline{[W^\ell]}$) behaves equivalently (split-
577 MNIST $92.2 \pm 0.4 / 95.2 \pm 0.3$, six seeds; CIFAR-10 $28.1 \pm 0.9 / 31.2 \pm 0.8$; feature CIFAR-10
578 $40.5 \pm 0.6 / 43.8 \pm 0.2$; CIFAR-100 $3.7 \pm 0.4 / 4.3 \pm 0.2$). The drive estimate must see only unamplified
579 waking updates: inflated one-hot targets and the offline night’s $3\times$ -gain batches push the EMA to
580 its clip and erase day learning (night $\times$ controller 60.2 ± 2.8 against 90.9 ± 3.2 for the night alone),
581 whereas anchor and bout gate compose cleanly (bout + controller 96.0 ± 0.3 static). The controller
582 does not subsume λ ’s role as a capacity regulariser against overfitting on small tabular problems,
583 information no drive ratio contains. On CIFAR-100 (ten tasks of ten classes; every method at floor,
584 chance 1%, BP+ER 9.9 ± 0.5 with twice the local learner’s forgetting) depth is a net loss that skips
585 only half refund (four hidden layers $1.9 / 1.6 \rightarrow 3.8 / 3.0$ with skips against $5.7 / 6.0$ at two), while
586 width at the short chain length lifts the local record to $6.7 \pm 0.3 / 7.5 \pm 0.2$ (1024–512; a $5\times$ buffer
587 adds nothing): the thin-signal regime wants capacity, not depth, and the static-axis wall (6–7.5%
588 under every substrate change) says the front-end is the binding constraint there. All runs reproduce
589 bitwise at a fixed CPU thread count; a different thread count changes the reduction order, which
590 k -WTA’s discontinuity amplifies into a different trajectory.

591 K Held-out replication of the headline orderings

592 Every configuration choice made during development — some four hundred sequential-axis configu-
593 rations — was made on the official test sets. As a check on the orderings that carry the claims, every
594 configuration was frozen and the headline configurations were re-run with a tenth of the training
595 set held out as the evaluation set (stratified, fixed split; training data during development, never
596 consulted by any selection, never trained on here), six seeds each (Table 5; the CIFAR-10 check is in
597 Table 7). Held-out numbers sit 0.9 points below test on average, the sequential-axis ranking correlates
598 at Spearman $\rho = 0.97$, and the headline MNIST orderings survive: always-on rotation and the bout
599 gate lead ($91.6 \pm 0.3 / 91.8 \pm 0.3$ vs. the night’s 89.0 ± 3.5 , BP+ER’s 88.8 ± 0.3 and unmasked ER’s
600 85.4 ± 3.5 ; paired bootstrap margins $+2.6$ [$+0.2, +5.2$], $+2.8$ [$+2.5, +3.2$] and $+6.3$ [$+3.8, +9.1$]
601 over six seed pairs); the silent mask trails rotation; the five-layer skip system holds the two-layer
602 level on both axes ($91.0 / 95.8$). One claim does not replicate cleanly: the adaptive decay’s static-side
603 margin (96.1 vs. 94.7 on test) becomes 93.5 ± 4.3 vs. 94.1 ± 0.2 on the held-out split, where one
604 of six seeds collapsed — we report it as unconfirmed. On CIFAR-10 (Table 7, three seeds) every
605 ordering replicates: raw $28.9 > 25.7 > 24.8 > 15.4 \approx 16.3$ for rotation, night, BP+ER, unmasked
606 ER and no buffer; features $42.3 > 34.0 \gg 20.5$ with BP+ER at 42.8 (0.5 ahead, against 1.7 on test)
607 and the decomposition chain $43.7 / 43.5 / 40.4$. Each is a single fixed split.

Table 5: **Held-out replication.** Sequential / static accuracy on the official test set (all available seeds, three to six) and on a held-out tenth of the training set (six seeds).

	sequential		static	
	test	held-out	test	held-out
always-on rotation (default)	92.7 $\pm$ 0.3	91.6 $\pm$ 0.3	94.7 $\pm$ 0.3	94.1 $\pm$ 0.2
bout-committed gate	92.3 $\pm$ 0.2	91.8 $\pm$ 0.3	—	—
adaptive λ (drive)	91.7 $\pm$ 0.2	90.6 $\pm$ 0.6	96.1 $\pm$ 0.2	93.5 $\pm$ 4.3
five hidden layers + skips + adaptive λ	91.5 $\pm$ 0.4	91.0 $\pm$ 0.3	96.2 $\pm$ 0.3	95.8 $\pm$ 0.3
silent mask (no rotation)	88.8 $\pm$ 0.8	88.0 $\pm$ 0.7	—	—
offline night, same substrate	90.9 $\pm$ 3.2	89.0 $\pm$ 3.5	—	—
unmasked interleaved ER	87.0 $\pm$ 1.1	85.4 $\pm$ 3.5	—	—
BP + ER	89.3 $\pm$ 0.8	88.8 $\pm$ 0.3	—	—
no buffer	19.6 $\pm$ 0.0	19.4 $\pm$ 0.0	—	—

Table 6: **Per-task accuracy matrix** of the headline configuration on split-MNIST (official test set, mean over six seeds): row = after training task i , column = accuracy on task j . Forgetting F , the mean of the diagonal minus the last row over tasks 1–4, is 5.3 ± 0.4 ; the drop is concentrated on tasks 2–3 and the first task is retained.

after task	T1	T2	T3	T4	T5
1	99.8				
2	99.3	98.4			
3	98.8	93.2	97.4		
4	98.6	91.4	92.9	96.7	
5	97.9	89.7	90.6	93.0	91.8

Table 7: **Held-out replication on CIFAR-10.** Sequential accuracy on the official test set (three seeds) and on a held-out tenth of the training set (three seeds), configurations frozen before the check.

		test	held-out
raw	refractory local sleep	29.5 $\pm$ 0.8	28.9 $\pm$ 1.0
	offline night	26.2 $\pm$ 0.6	25.7 $\pm$ 1.2
	BP + ER	25.1 $\pm$ 1.5	24.8 $\pm$ 0.3
	unmasked interleaved ER	17.0 $\pm$ 6.1	15.4 $\pm$ 4.2
	no buffer	16.4 $\pm$ 0.2	16.3 $\pm$ 0.1
	refractory, static	27.6 $\pm$ 0.4	28.9 $\pm$ 1.0STA
	unmasked, static	29.5 $\pm$ 2.8	15.4 $\pm$ 4.2STA
features	refractory local sleep	41.5 $\pm$ 1.1	42.3 $\pm$ 0.6
	offline night	35.0 $\pm$ 0.5	34.0 $\pm$ 0.8
	unmasked interleaved ER	20.1 $\pm$ 7.3	20.5 $\pm$ 1.0
	BP + ER (best rate)	43.2 $\pm$ 1.0	42.8 $\pm$ 0.7
	BP + ER, wide	44.5 $\pm$ 1.0	43.7 $\pm$ 1.0
	+ 30% wiring	43.3 $\pm$ 0.9	43.5 $\pm$ 0.2
	+ k -WTA	40.5 $\pm$ 0.7	40.4 $\pm$ 1.5